

# OSA CON 25



## CLICKHOUSE® CHRONICLES: REAL-WORLD WAR ROOMS WITH HUMAN AND AI AGENTS

NOVEMBER 4 & 5, 2025  
ONLINE EVENT



Shivji Kumar Jha

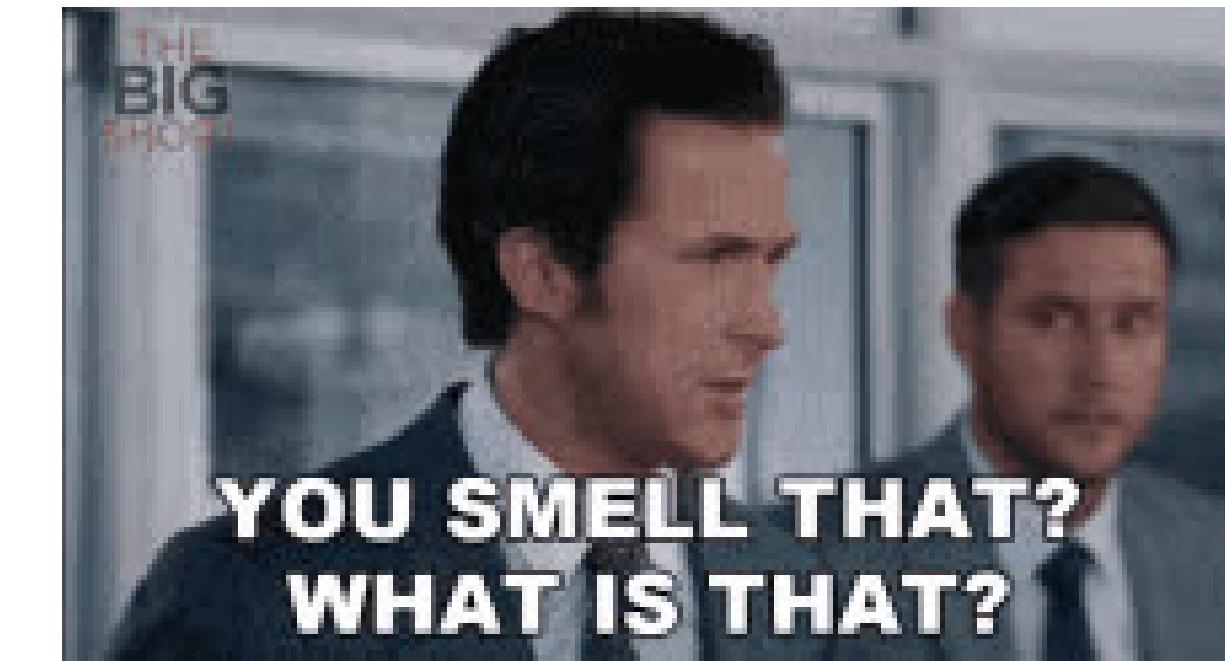
Staff Engineer @ Nutanix  
Data & Cloud geek



Anurag Pandey

Building Ai Copilot  
for Databases

# AI + L1-2 Clickhouse production problems



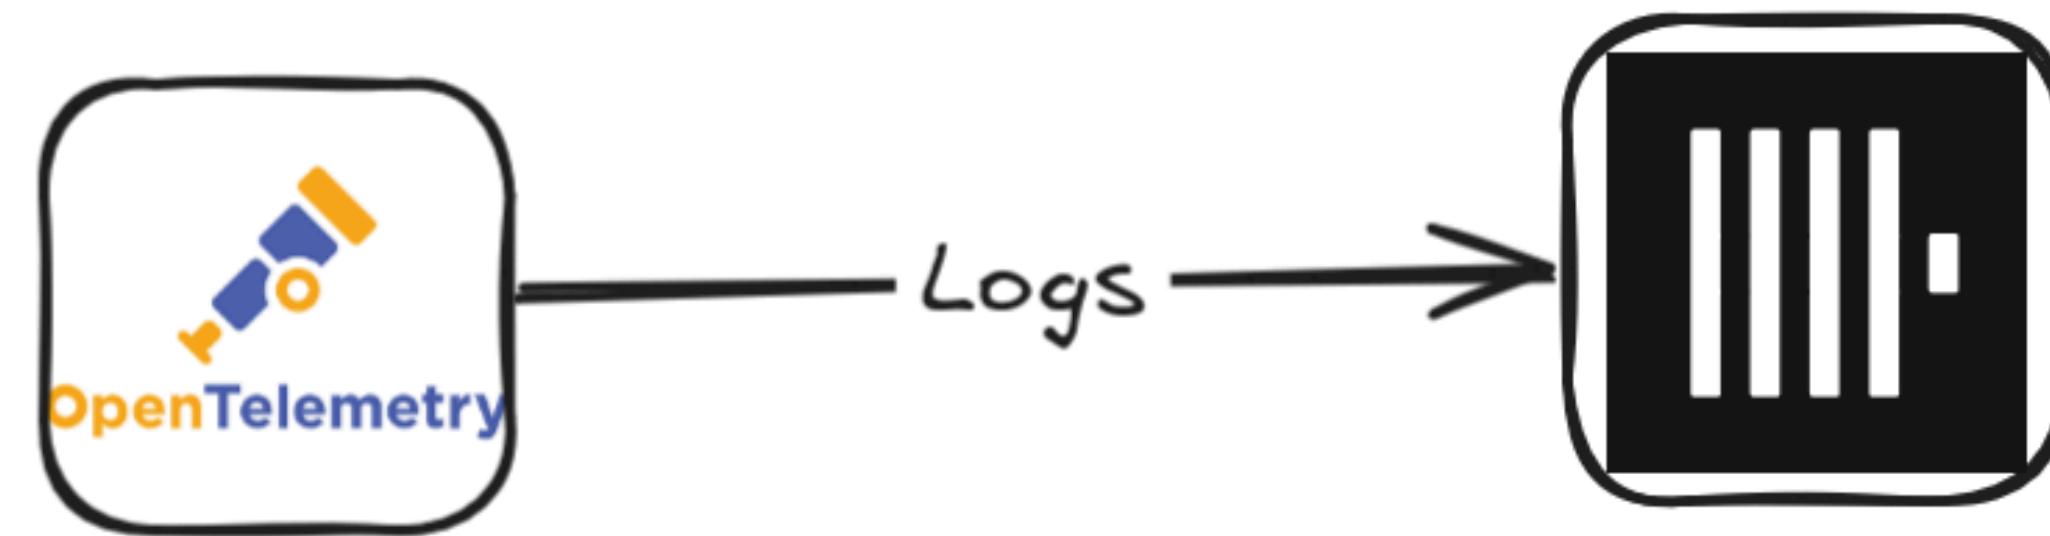
incerto

# S3 spikes 100x in cost

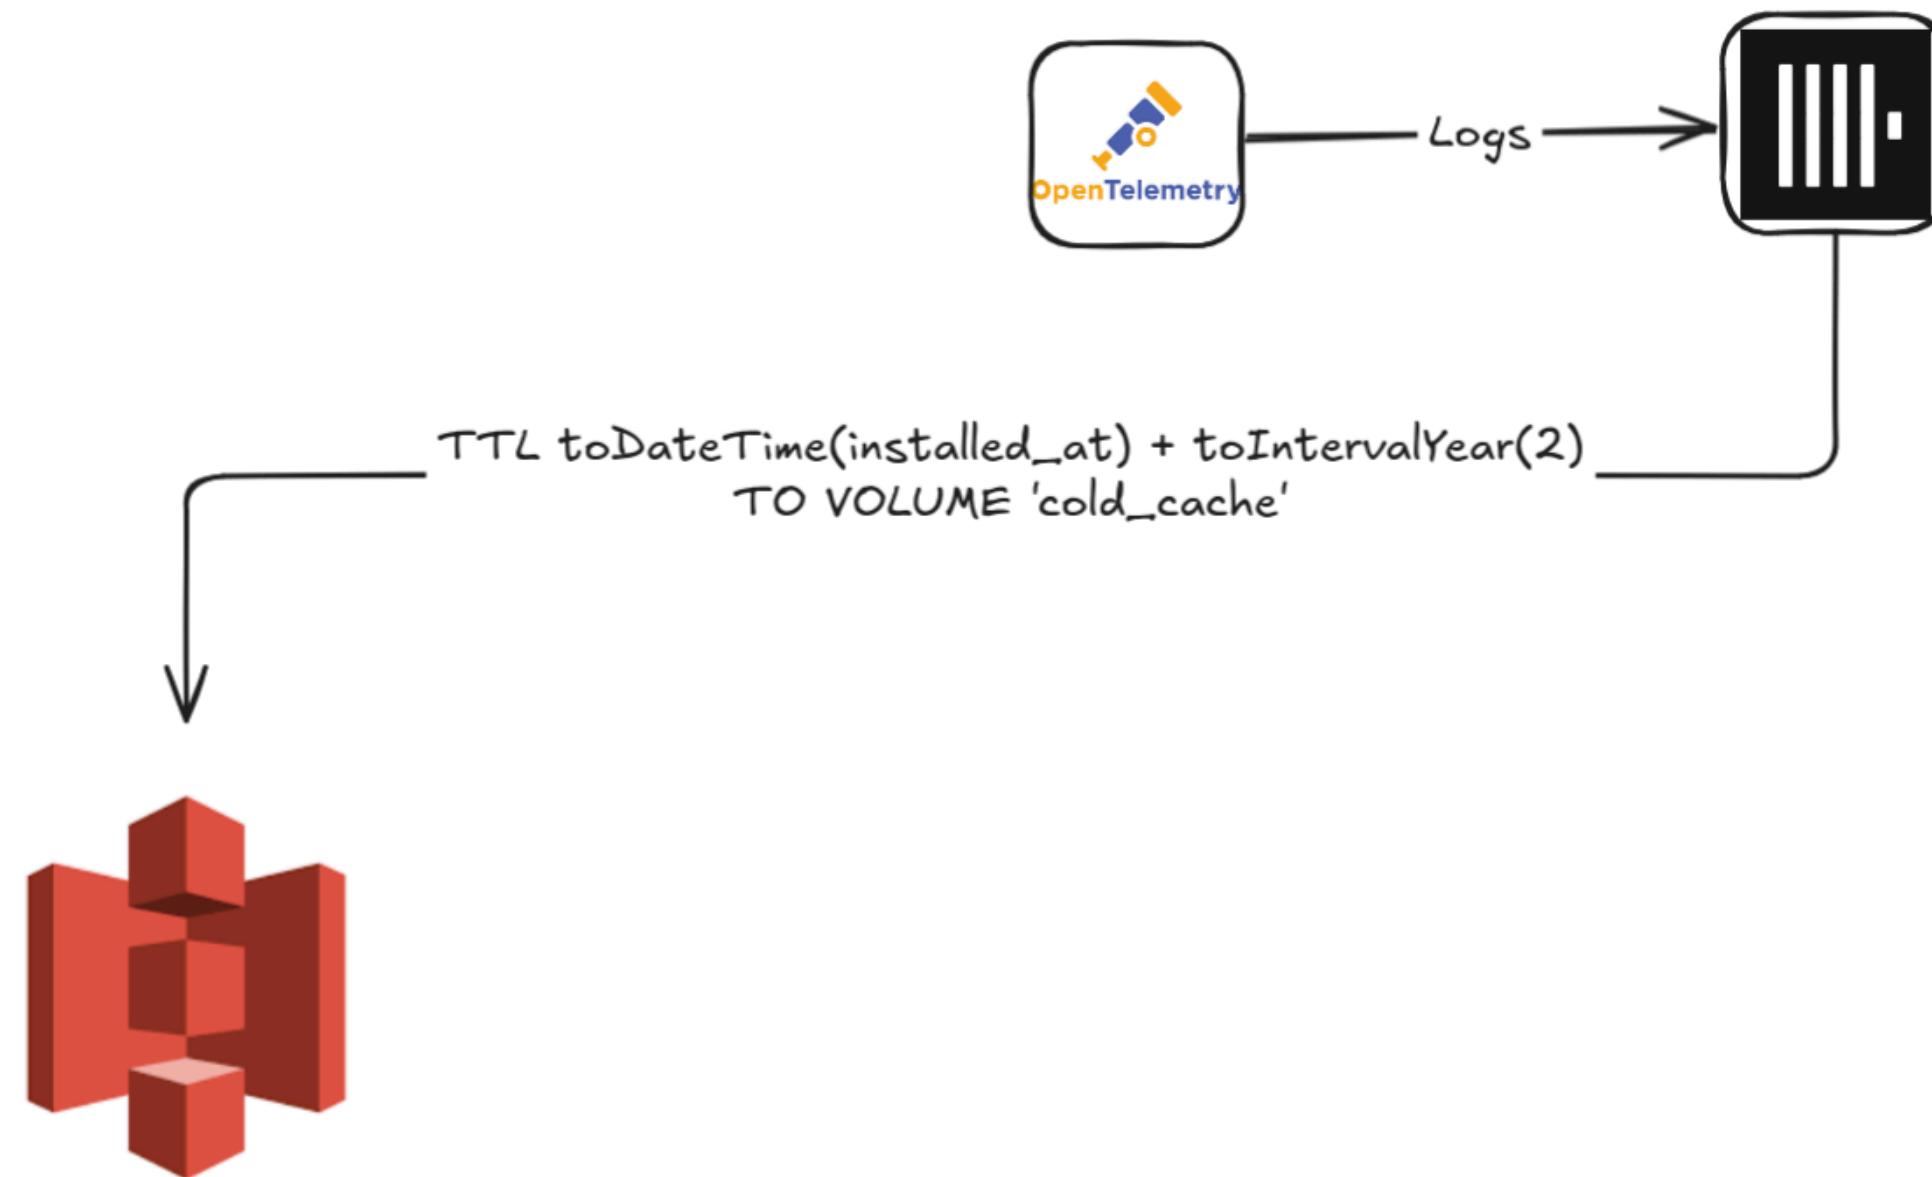


incerto

# Otel Logs being sent to Clickhouse

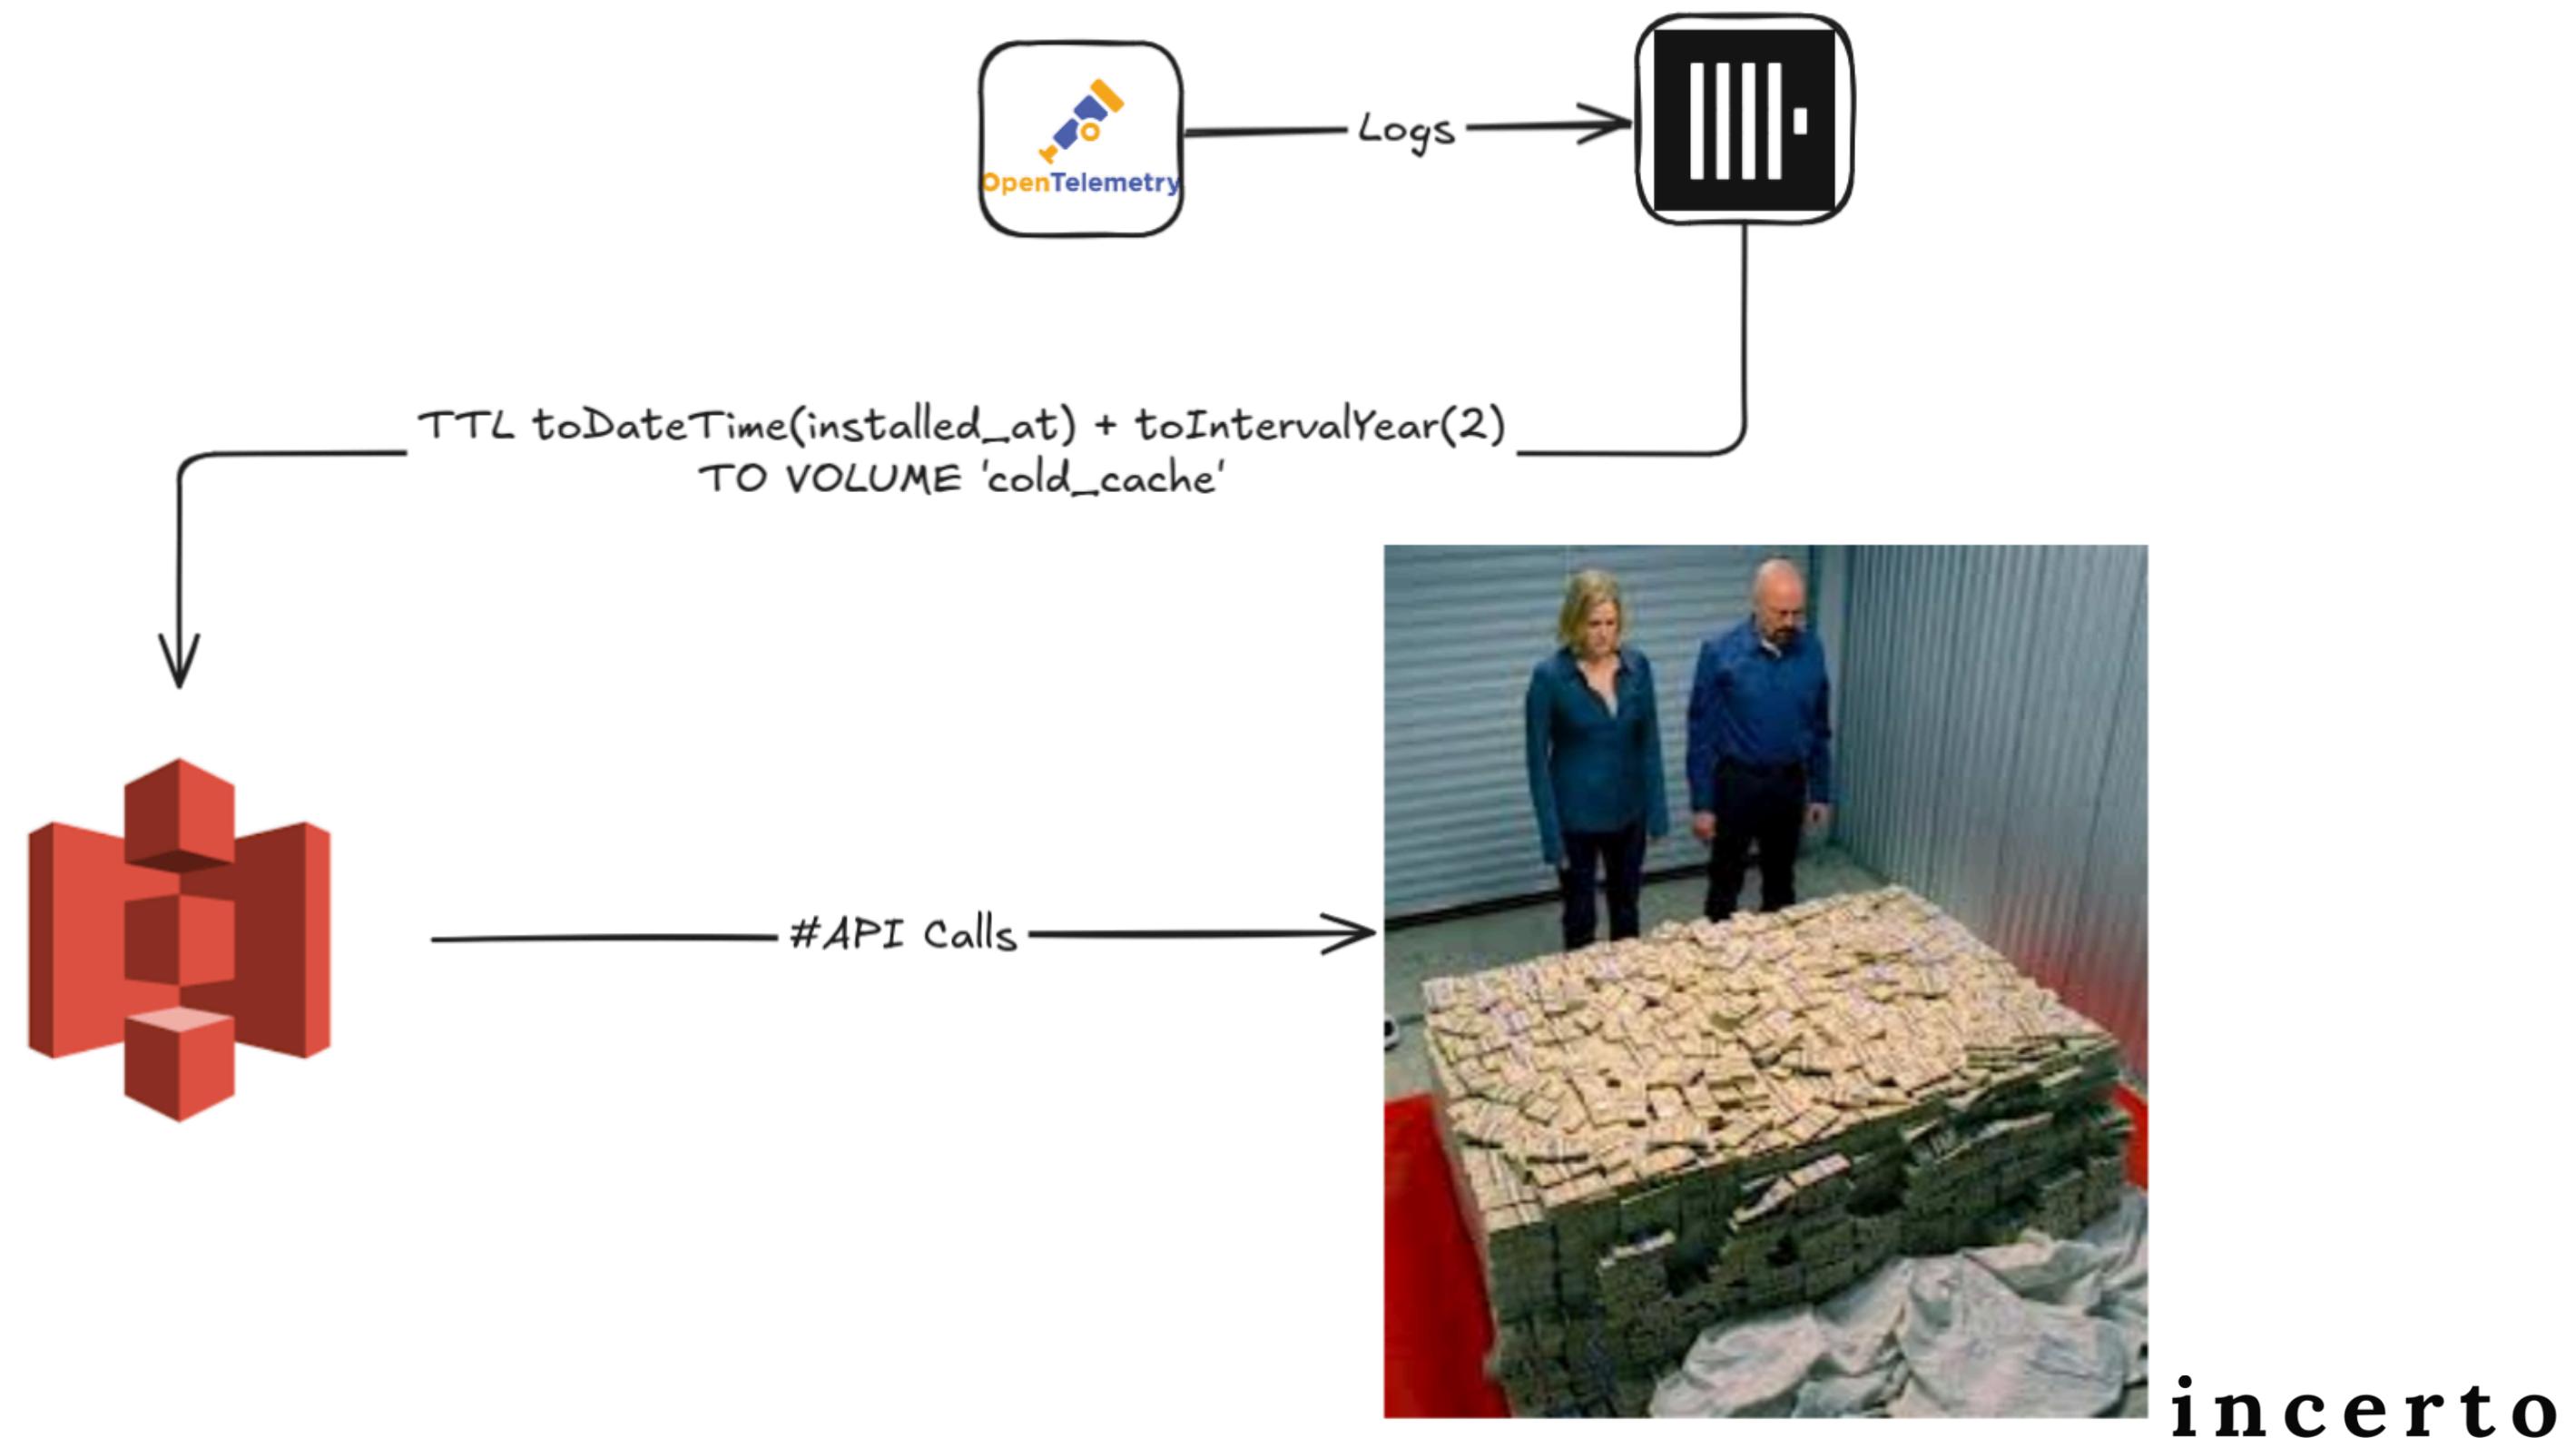


# 2yrs old logs moved to s3



incerto

# API calls = Money



Root Cause ?

1970-01-01

incerto

# How did AI help ?

- User : “My S3 cost has spiked up”
- AI : “Checks for the object\_storage table”
- AI: Finds that there are lot of API calls
- AI : Guesses that it could due to charge on API
- User does the rest of root cause finding.

incerto

# User queries failing



incerto

1. QPS was the same
2. All Worker UP
3. No unusual queries



Root Cause ?

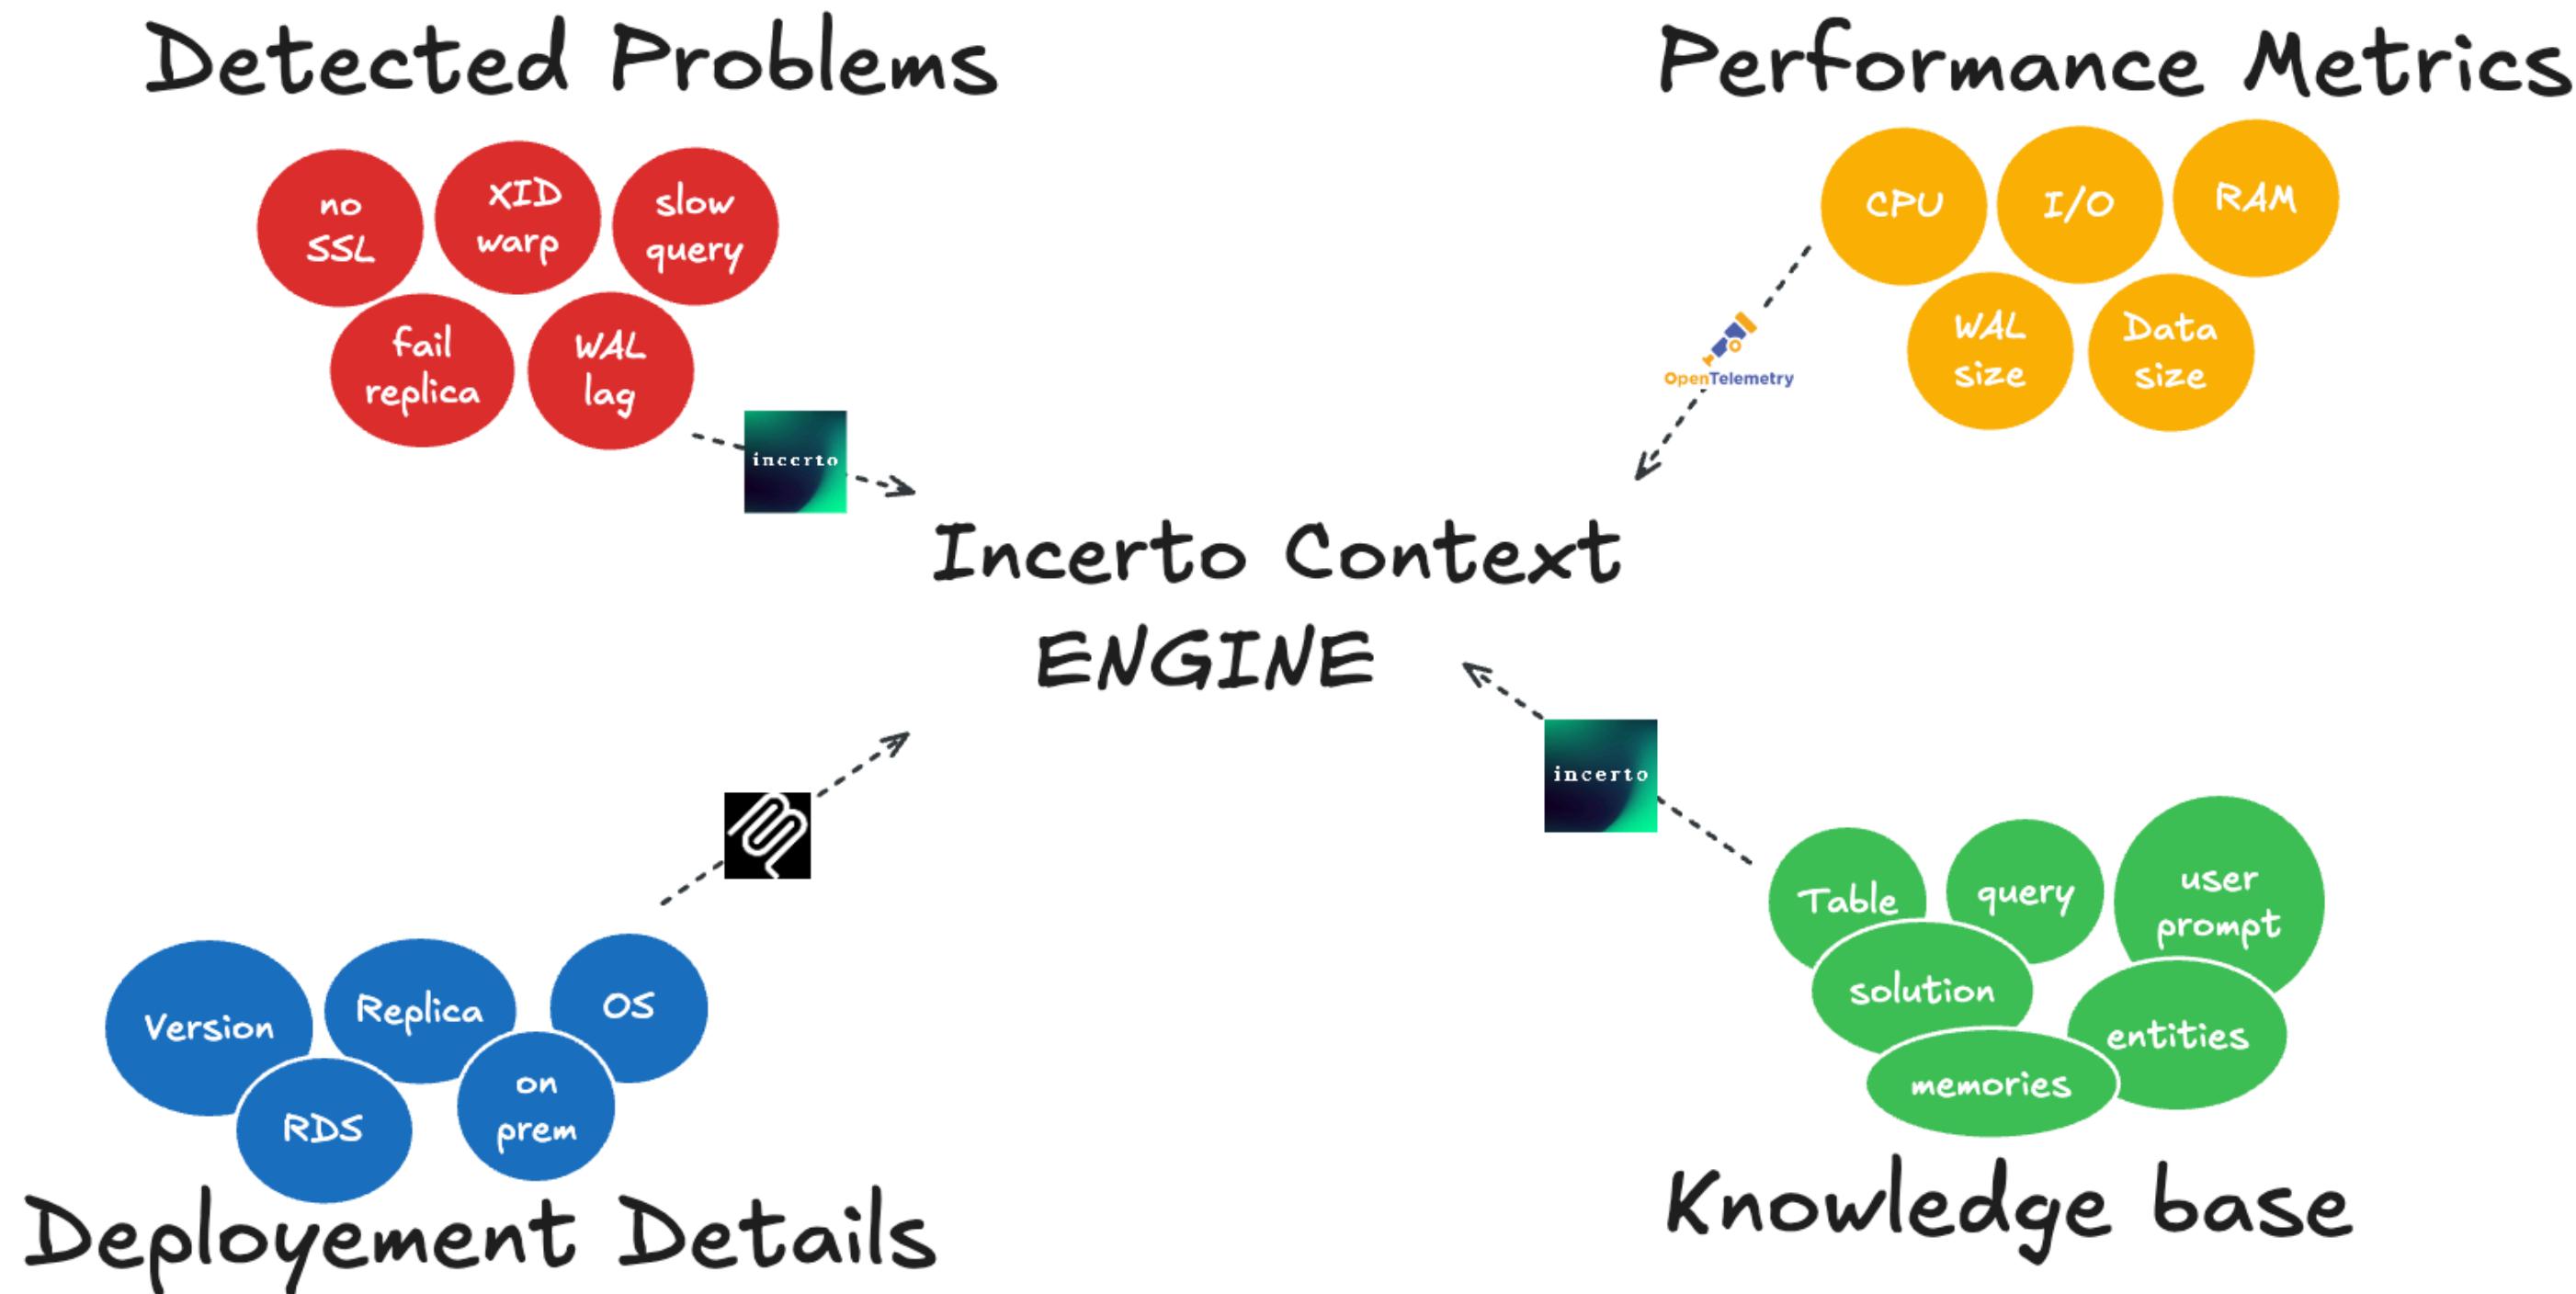
Huge Delete Mutations

incerto

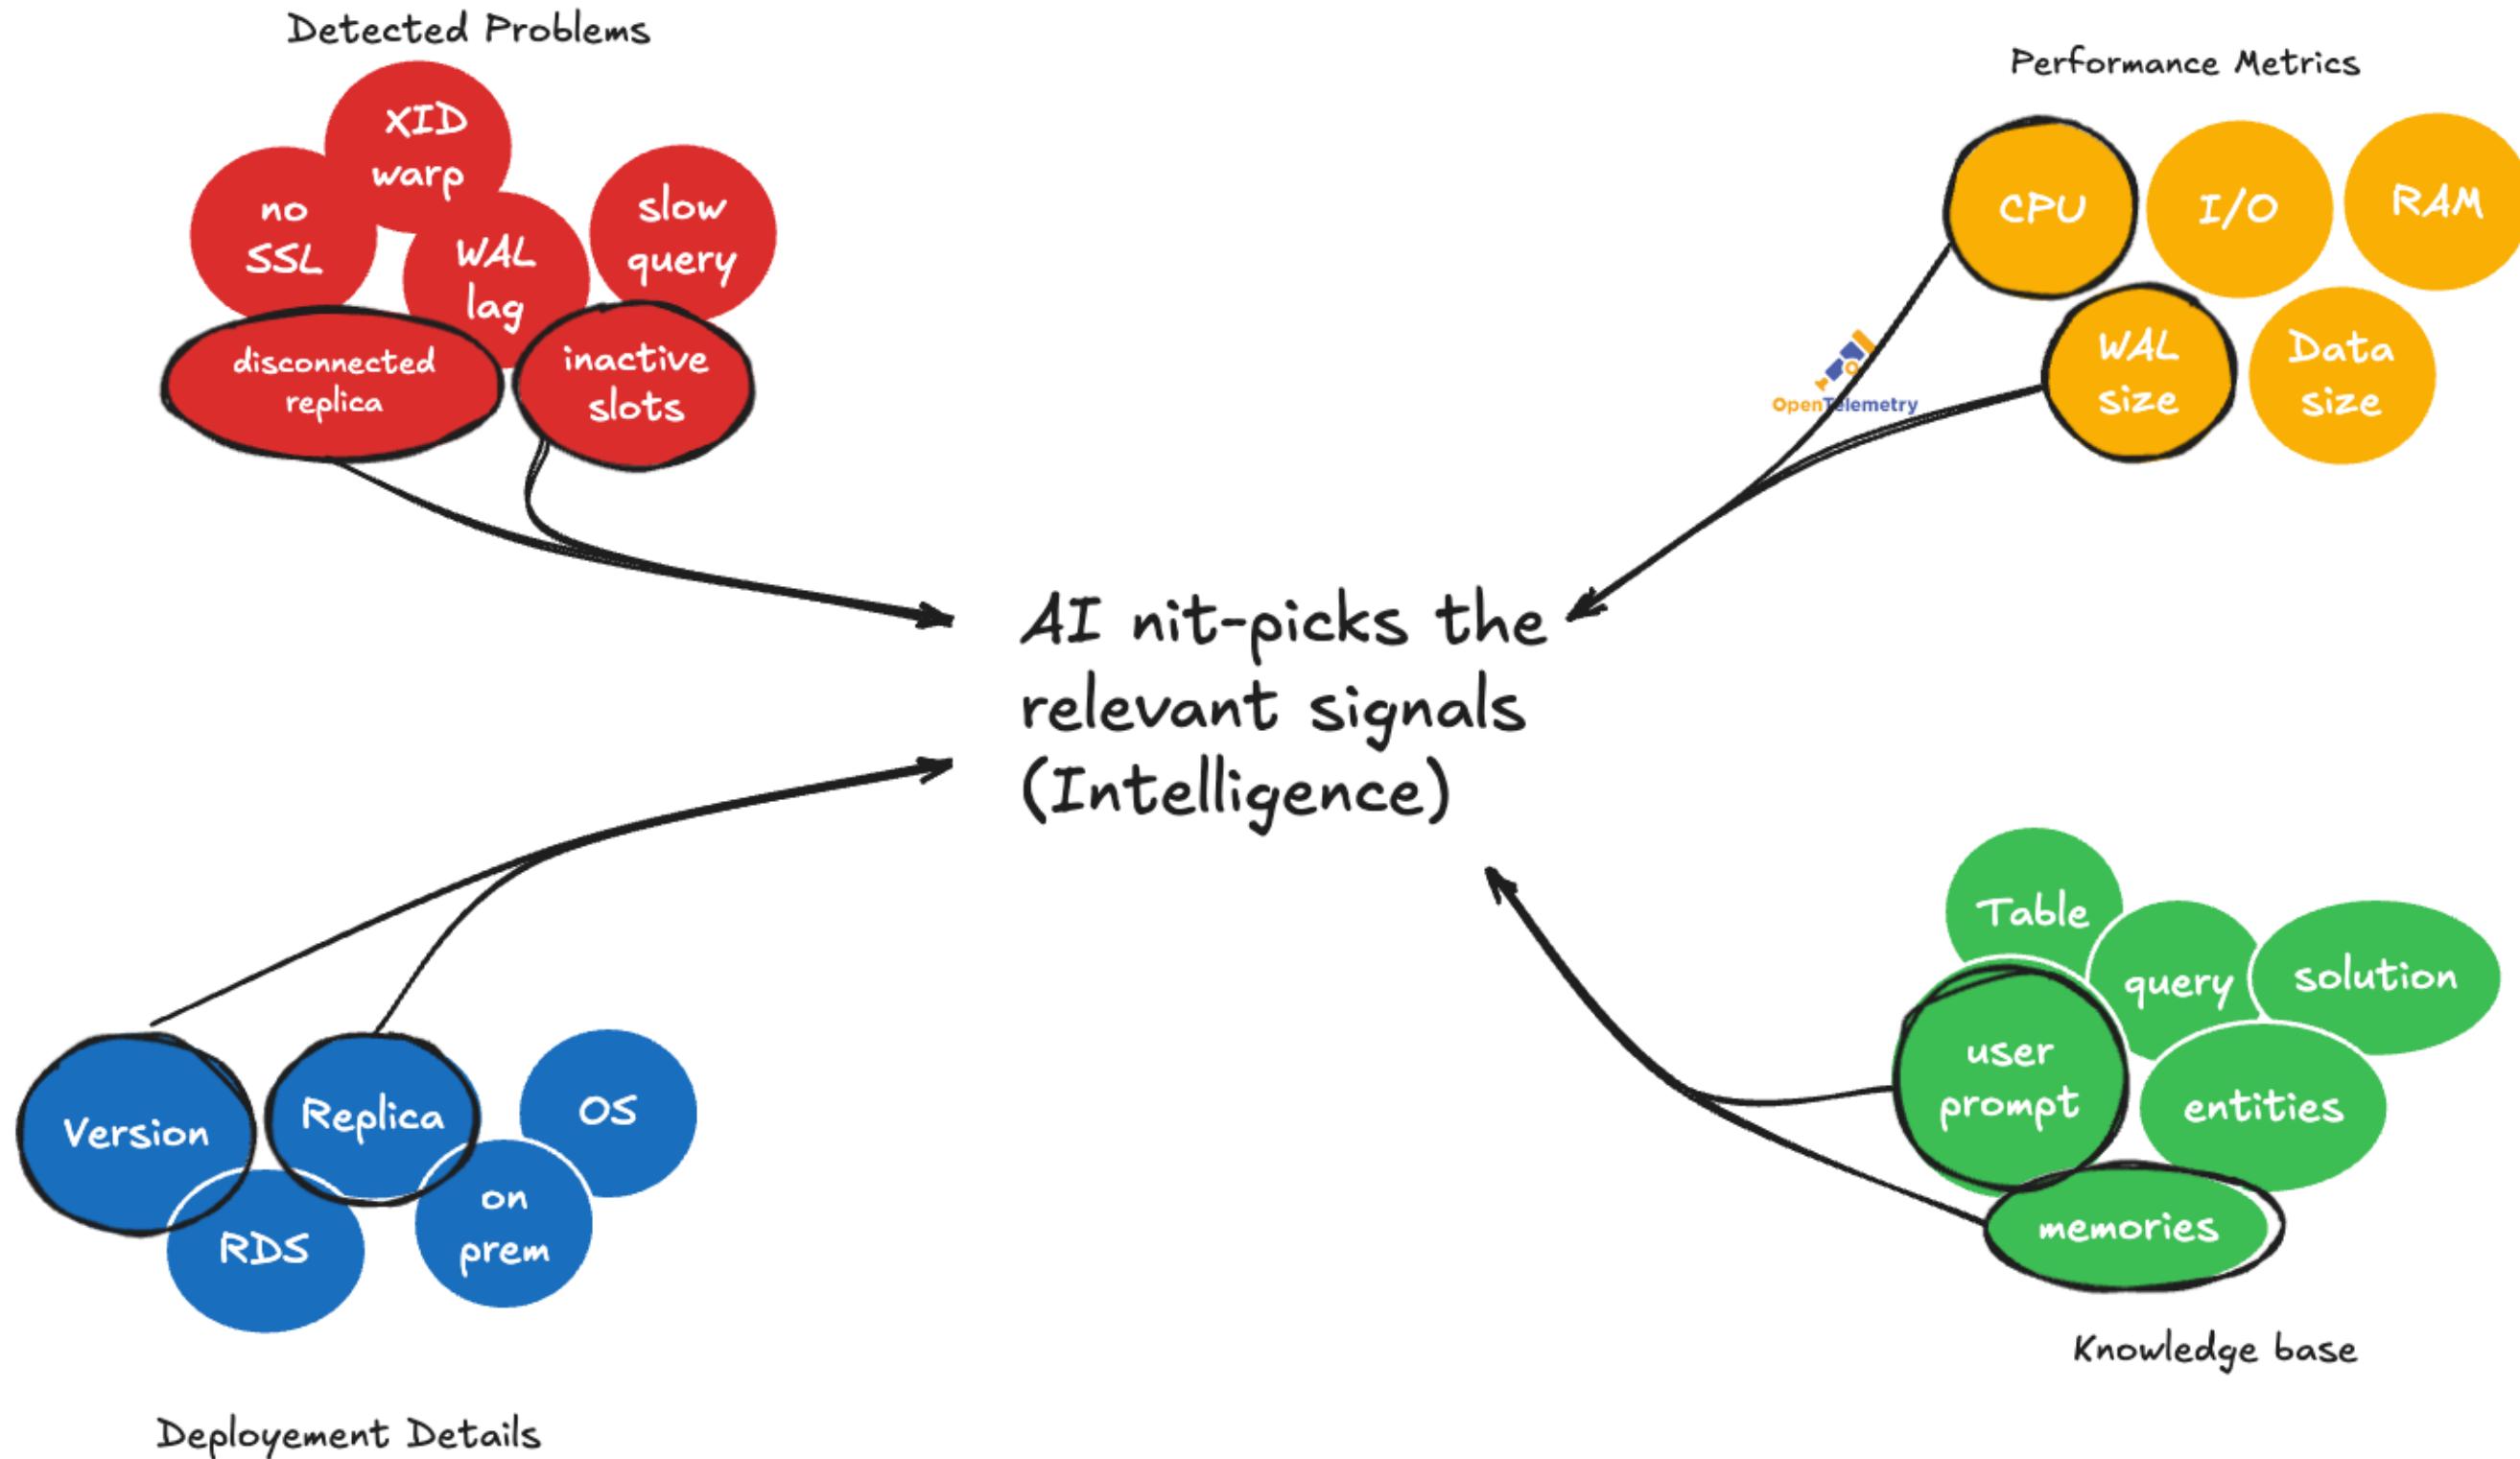
# How did AI help ?

- User : “I am seeing CPU spike”
- AI : Find the mutations running
- AI: Kills it at the users approval

# Context Engine

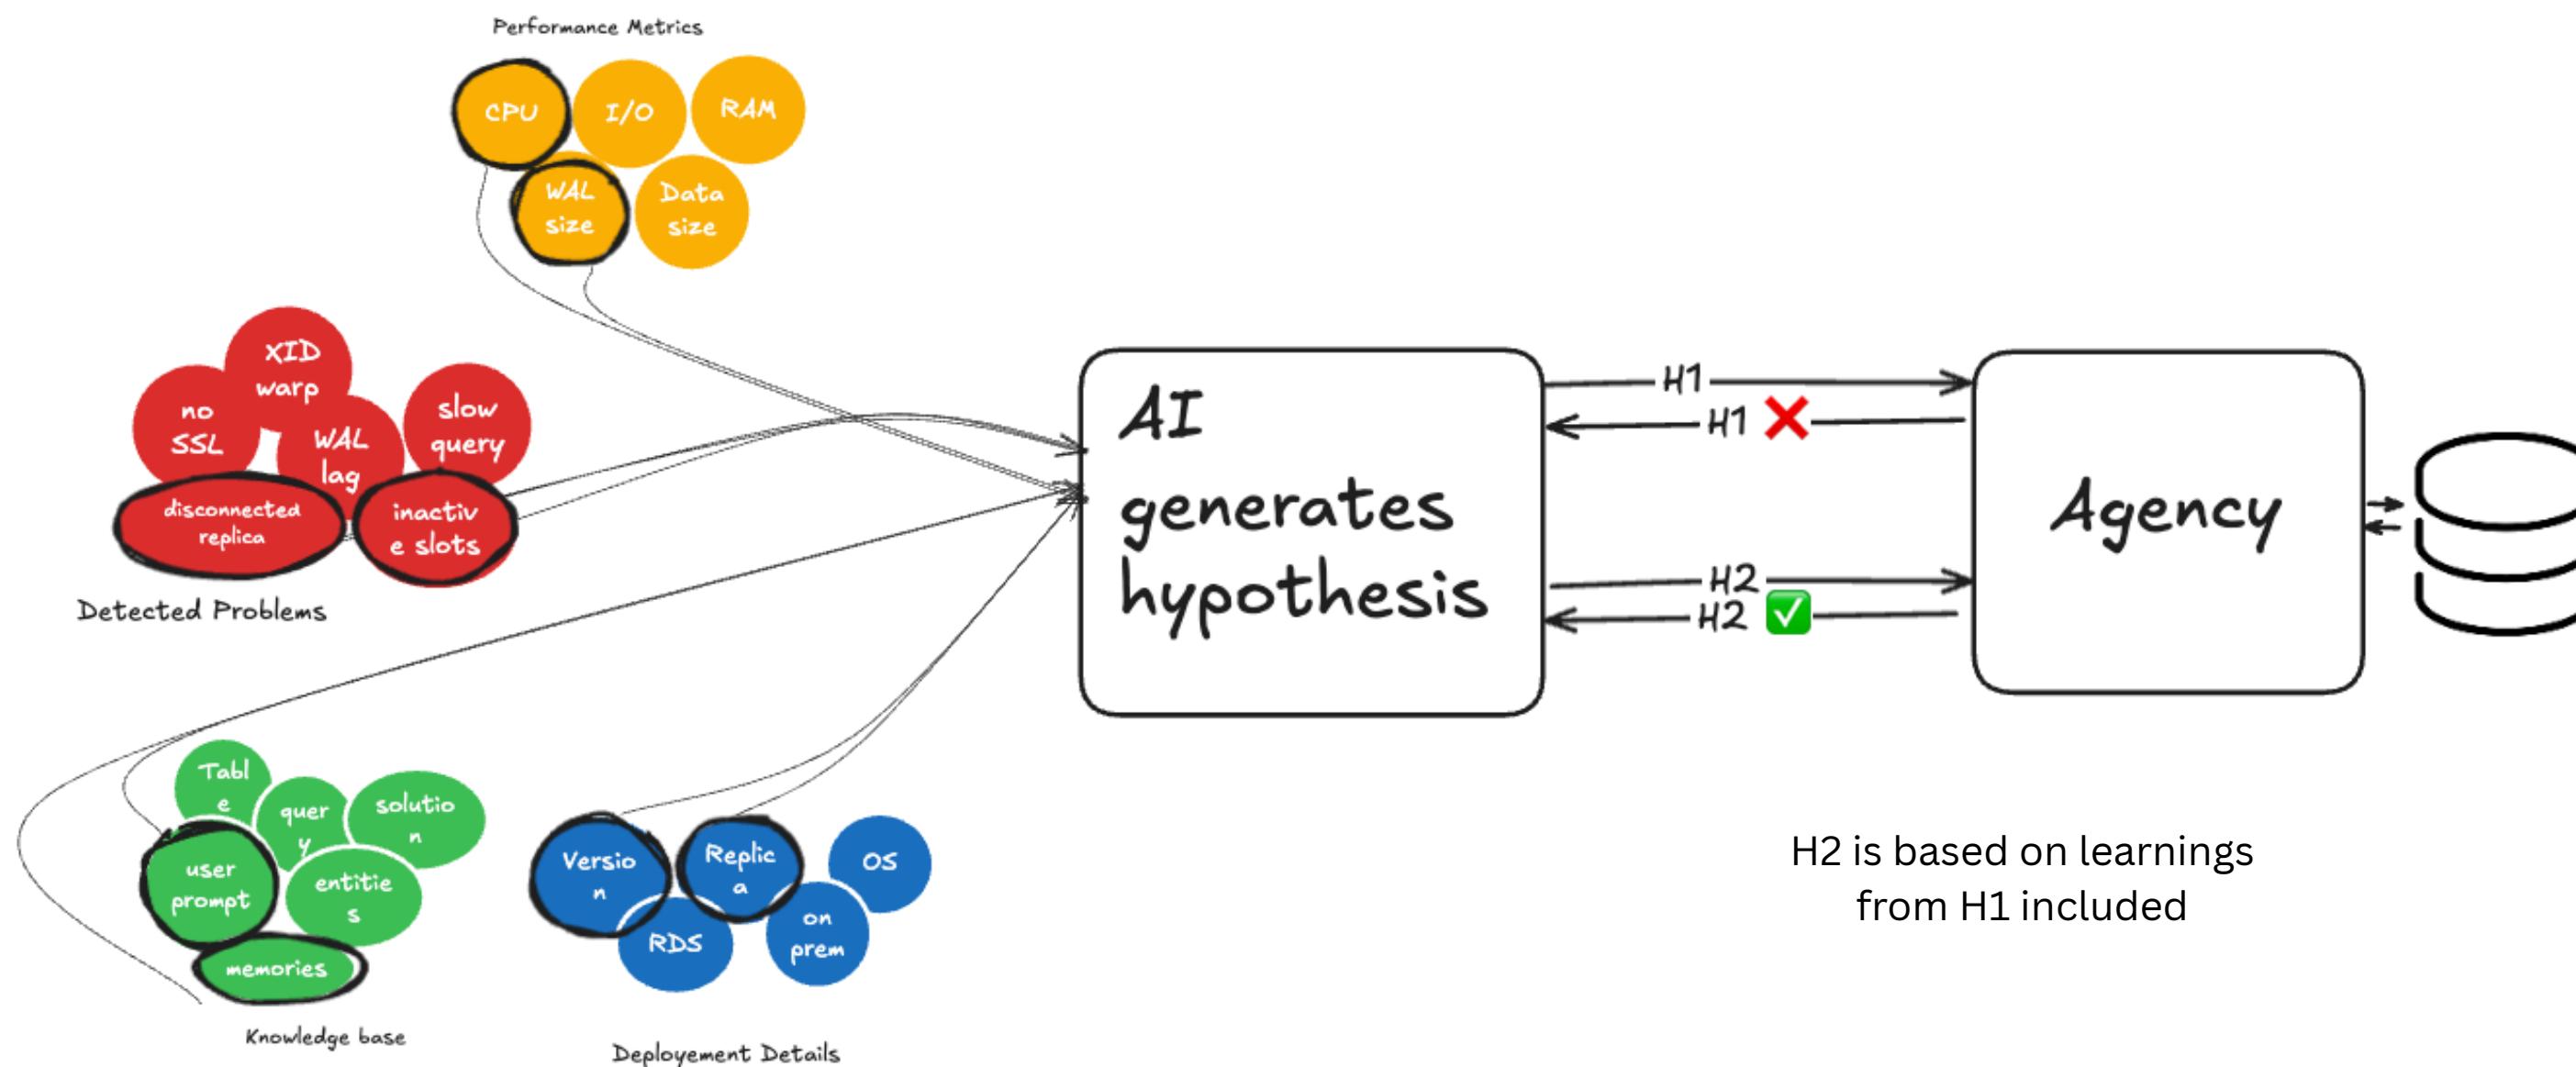


# AI picks up the right signals



# How do design AI Debugger

## AI generates & verifies hypothesis



incerto

Don't trust anyone who tells you  
**"AI can do everything!"**



Lets look at some human actions...

# Distributed systems is hard!

Resiliency depends on  
how well you  
test **Failure scenarios**



# Some Chaos tests:

1. Introduce OS Kernel panic
2. Artificial Network latency
3. Stop the io channels abruptly
4. Kill pod, VM, power off abruptly



# Clickhouse server failed to come up

**Code: 722. DB::Exception:** Waited job failed: Code: 696. DB::Exception: Load job 'startup table db1.table1' -> Code: 695. DB::Exception: Load job 'load table db1.table1' failed: Code: 231. DB::Exception: Suspiciously many (119 parts, 0.00 B in total) broken parts to remove while **maximum allowed broken parts count is 100**. You can change the maximum value with merge tree setting 'max\_suspicious\_broken\_parts' in <merge\_tree> configuration section or in table settings in .sql file (don't forget to return setting back to default value): Cannot attach table `db1`.`table1` from metadata file /var/lib/clickhouse/store/402/40286535-a77d-4d49-890cfea3bf2f2d2e/table1.sql from query

.....

**(TOO\_MANY\_UNEXPECTED\_DATA\_PARTS),. (ASYNC\_LOAD\_WAIT\_FAILED)** (version 24.8.8.17 (official build))

# `max_suspicious_broken_parts`

1. For the write throughput of modern hardware,  
default value of **100 is too low!**
2. We changed it to 100K
3. We prefer self-healing over data loss
  - a. Getting access to cluster is very hard
  - b. Also, its observability data



[https://clickhouse.com/docs/operations/settings/merge-tree-settings#max\\_suspicious\\_broken\\_parts](https://clickhouse.com/docs/operations/settings/merge-tree-settings#max_suspicious_broken_parts)

## Human war story - 2

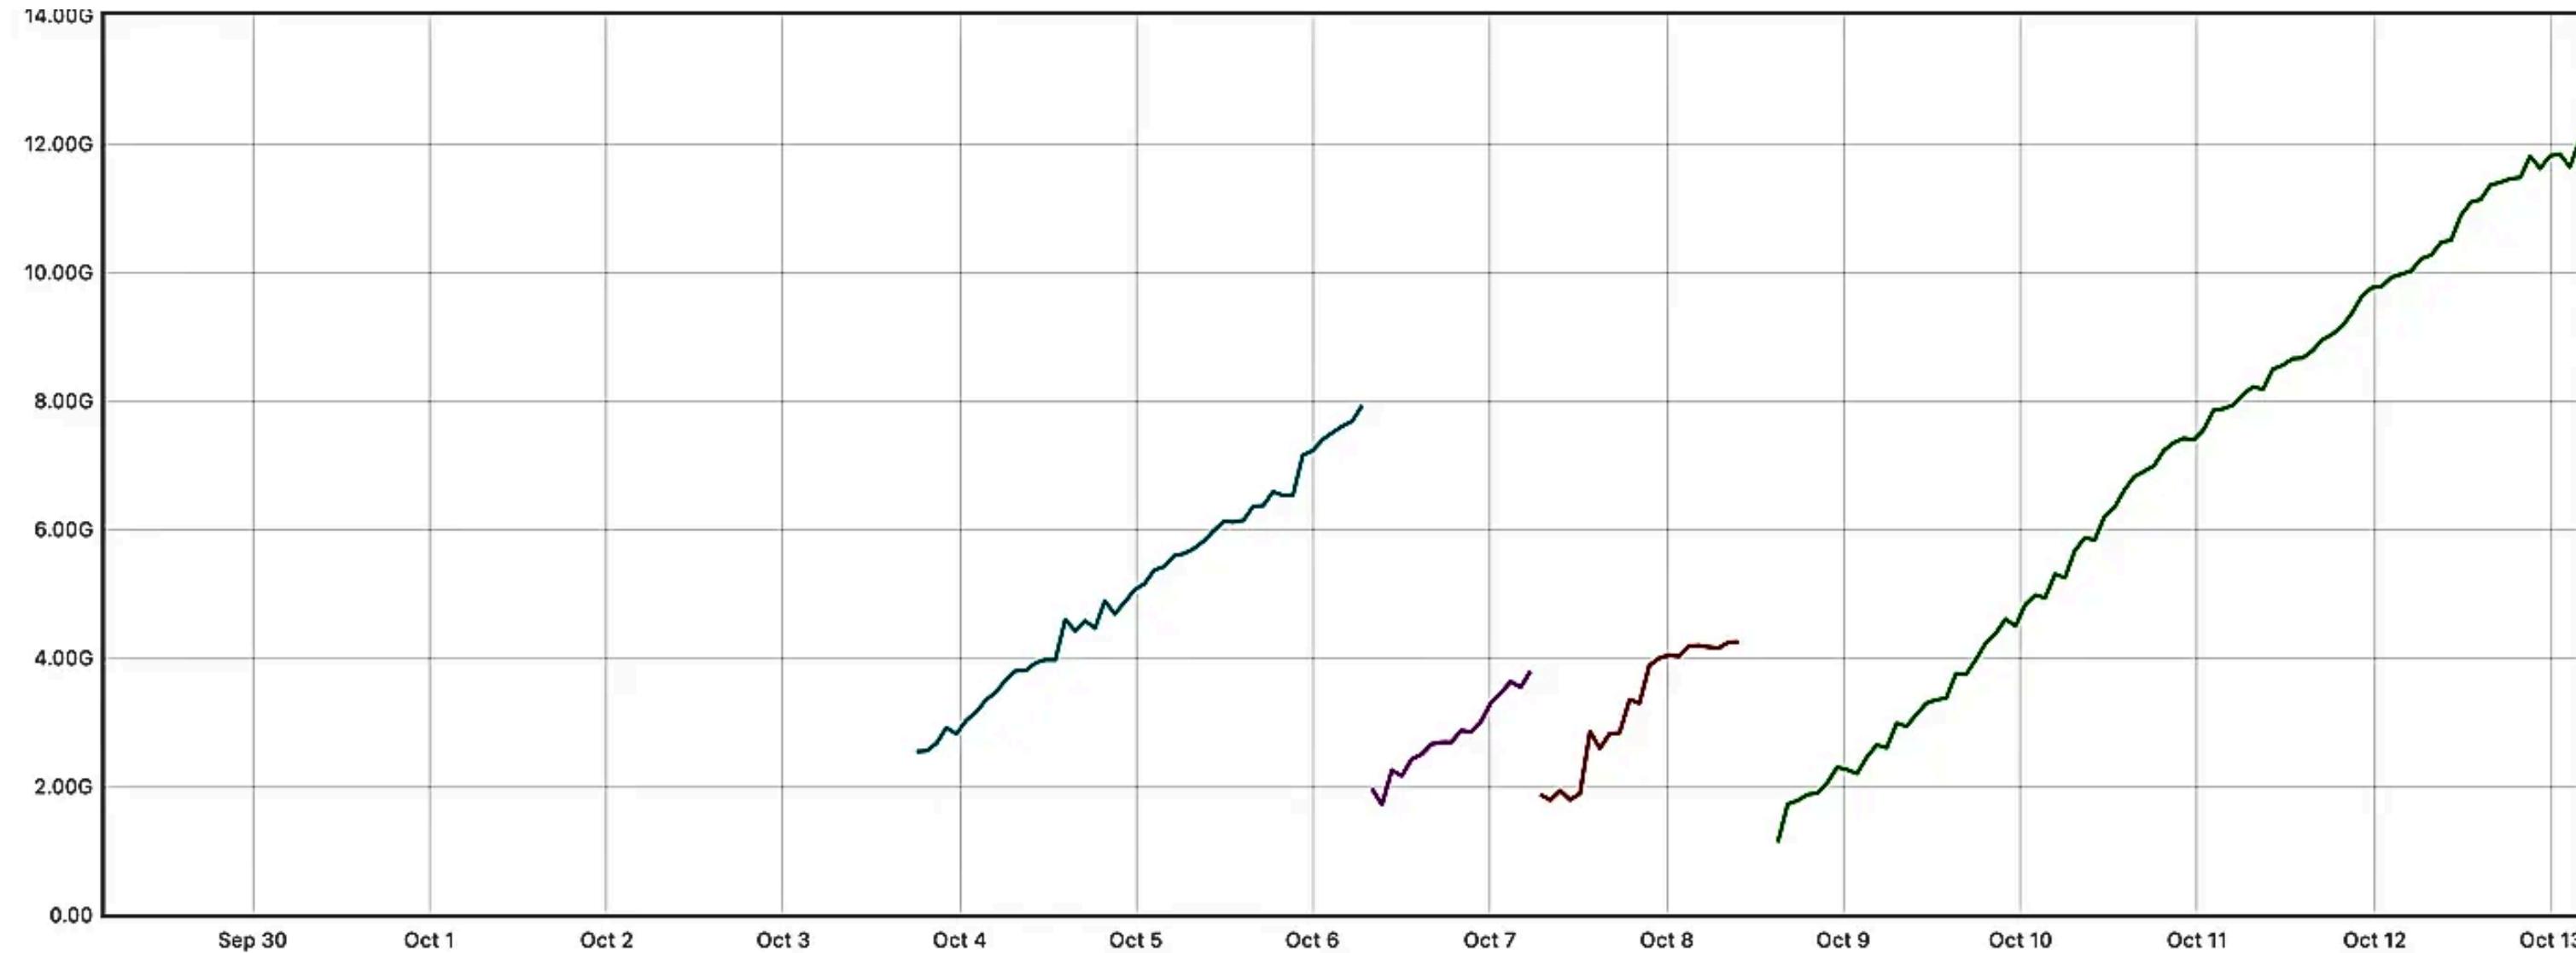
Queries began failing with  
**“Memory limit exceeded” errors**  
on ClickHouse v25.3.5.42

Every 5 to 6 days, clickhouse restarts 😢

thanks to the kubernetes gods



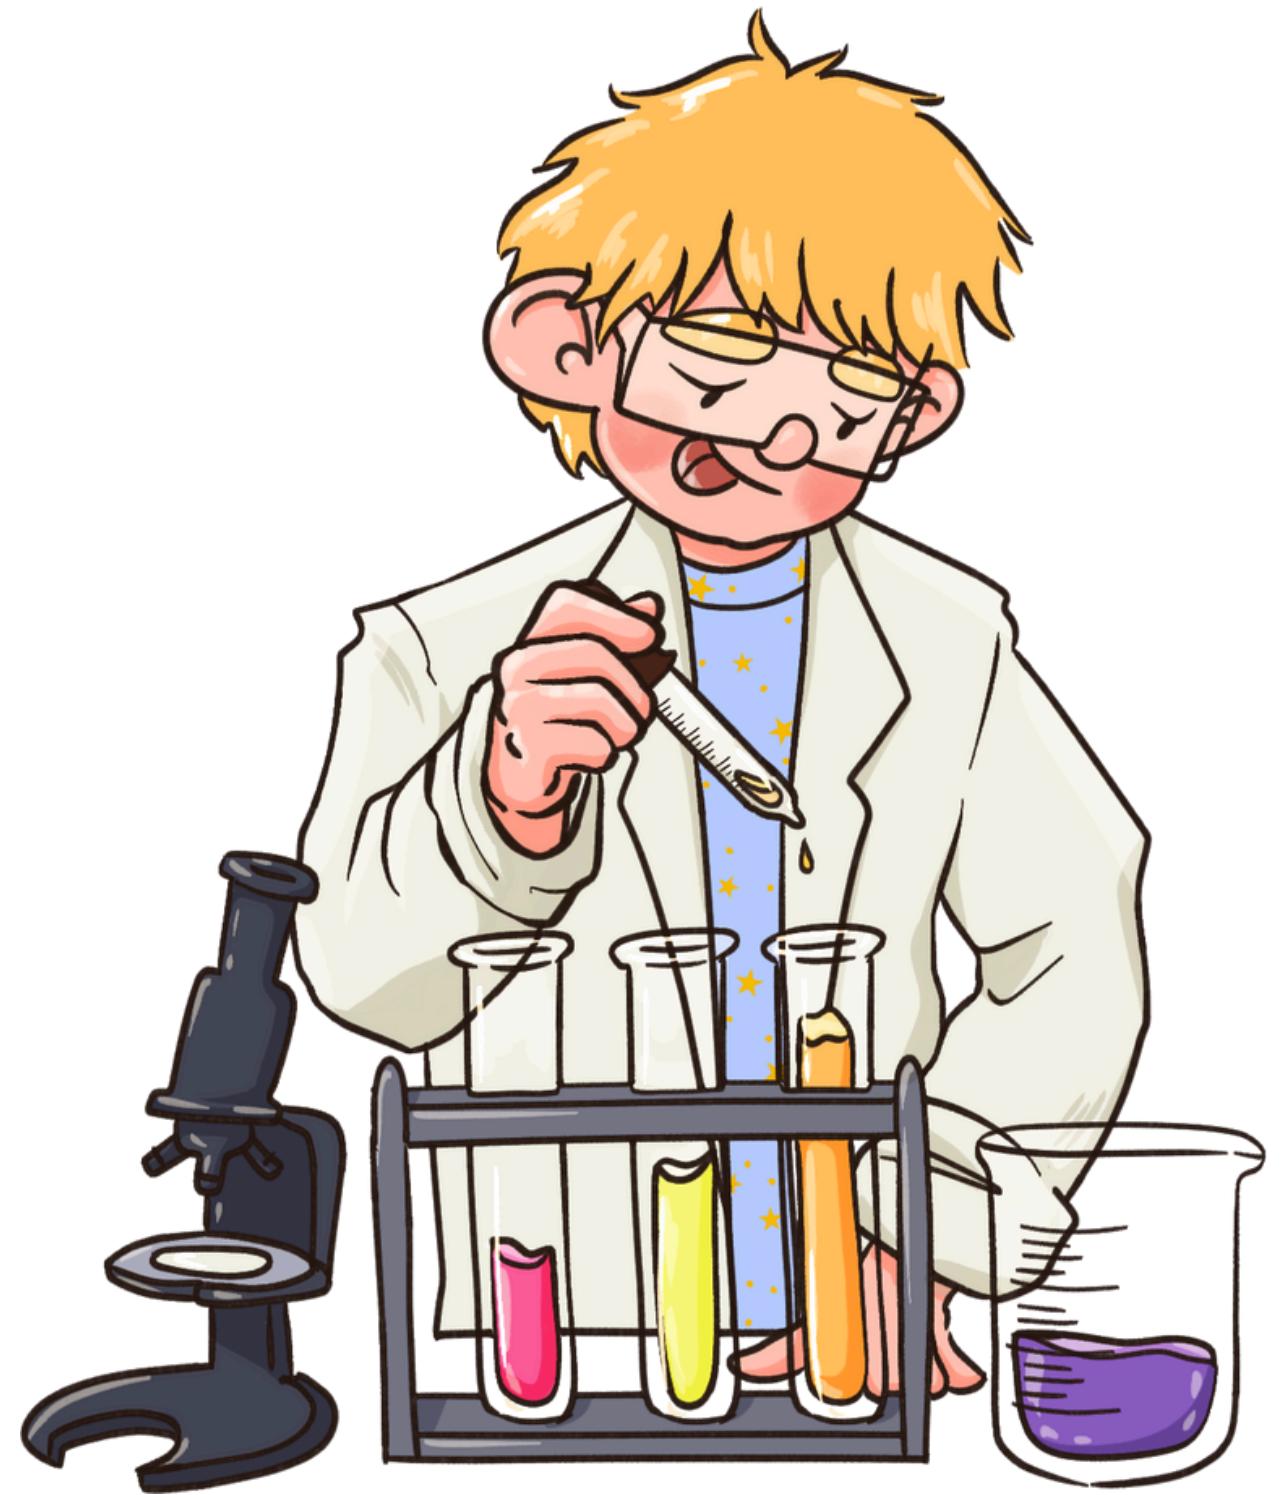
# The dreaded graph no dev wants to see (ouch!)



container\_memory\_working\_set\_bytes{container=~"clickhouse-pod"}

# Started with the usual suspects

1. check the query load
2. Investigate background merges & caches
3. Suspecting the Page cache
4. Nothing helped



Finally, accept the worst nightmare &  
profile **Jemalloc** with *jeprof*

# Fix server and merge upstream - thanks to the OSS gods!

**GLOBAL IN causes Memory Leak #88615**

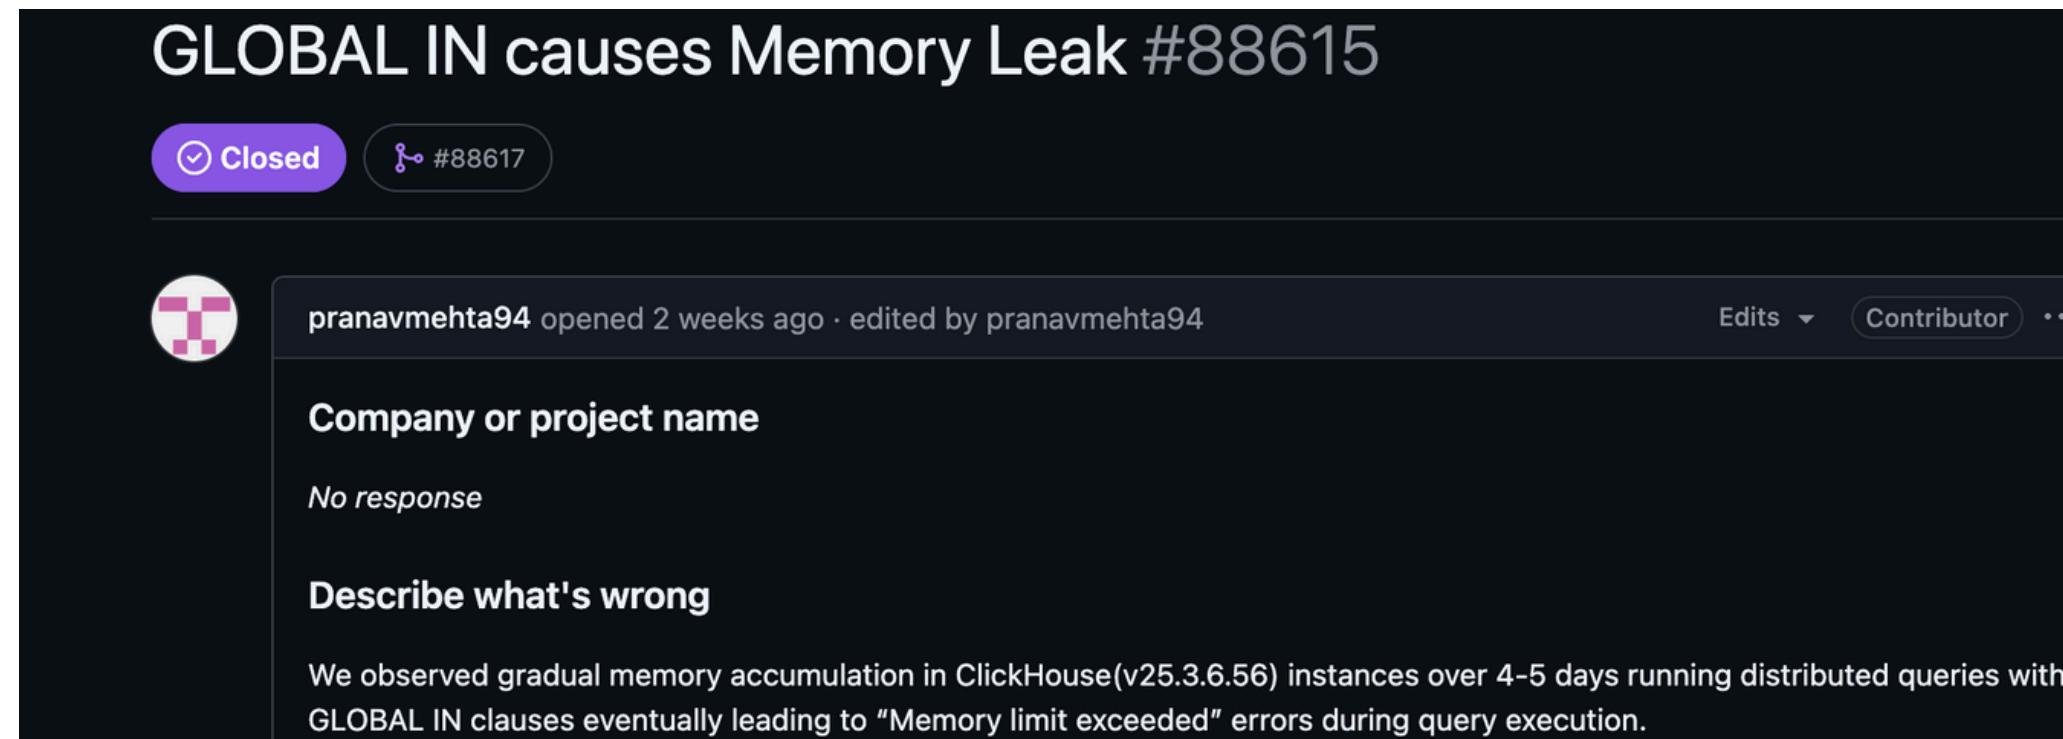
**Closed** #88617

pranavmehta94 opened 2 weeks ago · edited by pranavmehta94

Edits · Contributor · ...

**Company or project name**  
*No response*

**Describe what's wrong**  
  
We observed gradual memory accumulation in ClickHouse(v25.3.6.56) instances over 4-5 days running distributed queries with GLOBAL IN clauses eventually leading to "Memory limit exceeded" errors during query execution.



<https://github.com/ClickHouse/ClickHouse/issues/88615>

**cleanup temporary table entry from snapshot\_detached\_tables during table drop #88617**

**Merged** nickitat merged 2 commits into ClickHouse:master from pranavmehta94:pranav/global-in-fix 2 weeks ago

Conversation 2 · Commits 2 · Checks 119 · Files changed 1

pranavmehta94 commented 2 weeks ago · edited

Temporary tables are auto-created when using distributed queries with GLOBAL IN. These tables use the Memory Database engine and are passed to remote servers during distributed query processing.  
In DatabaseMemory::dropTable, when a table is dropped, the cleanup steps are as follows:

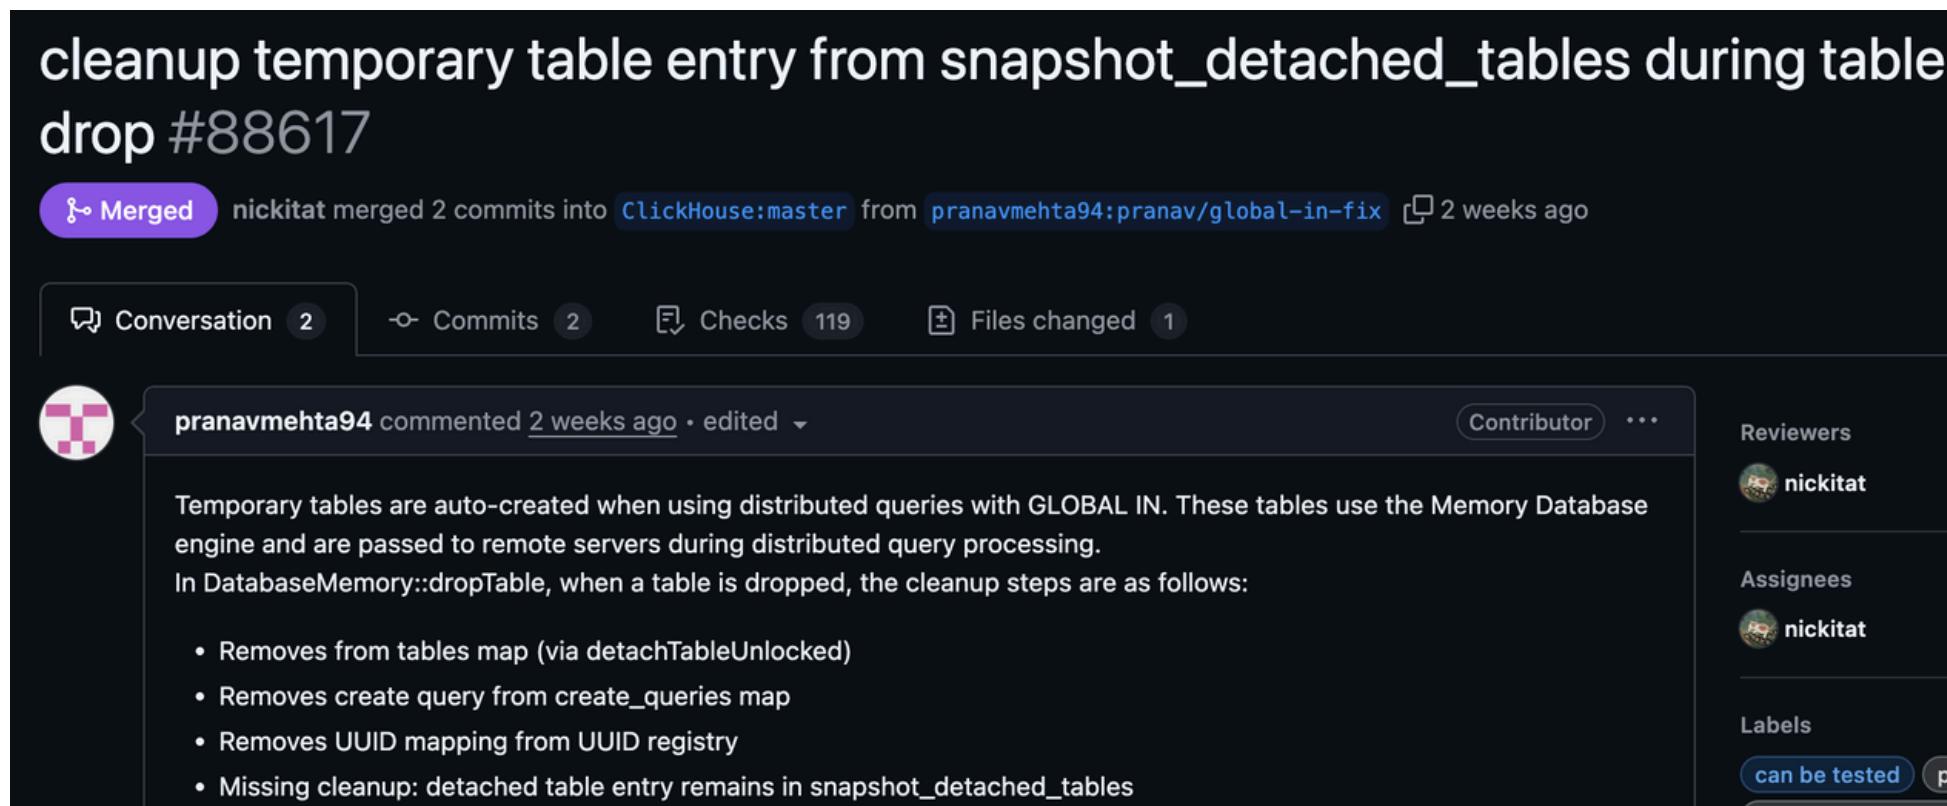
- Removes from tables map (via detachTableUnlocked)
- Removes create query from create\_queries map
- Removes UUID mapping from UUID registry
- Missing cleanup: detached table entry remains in snapshot\_detached\_tables

Contributor · ...

Reviewers  
nickitat

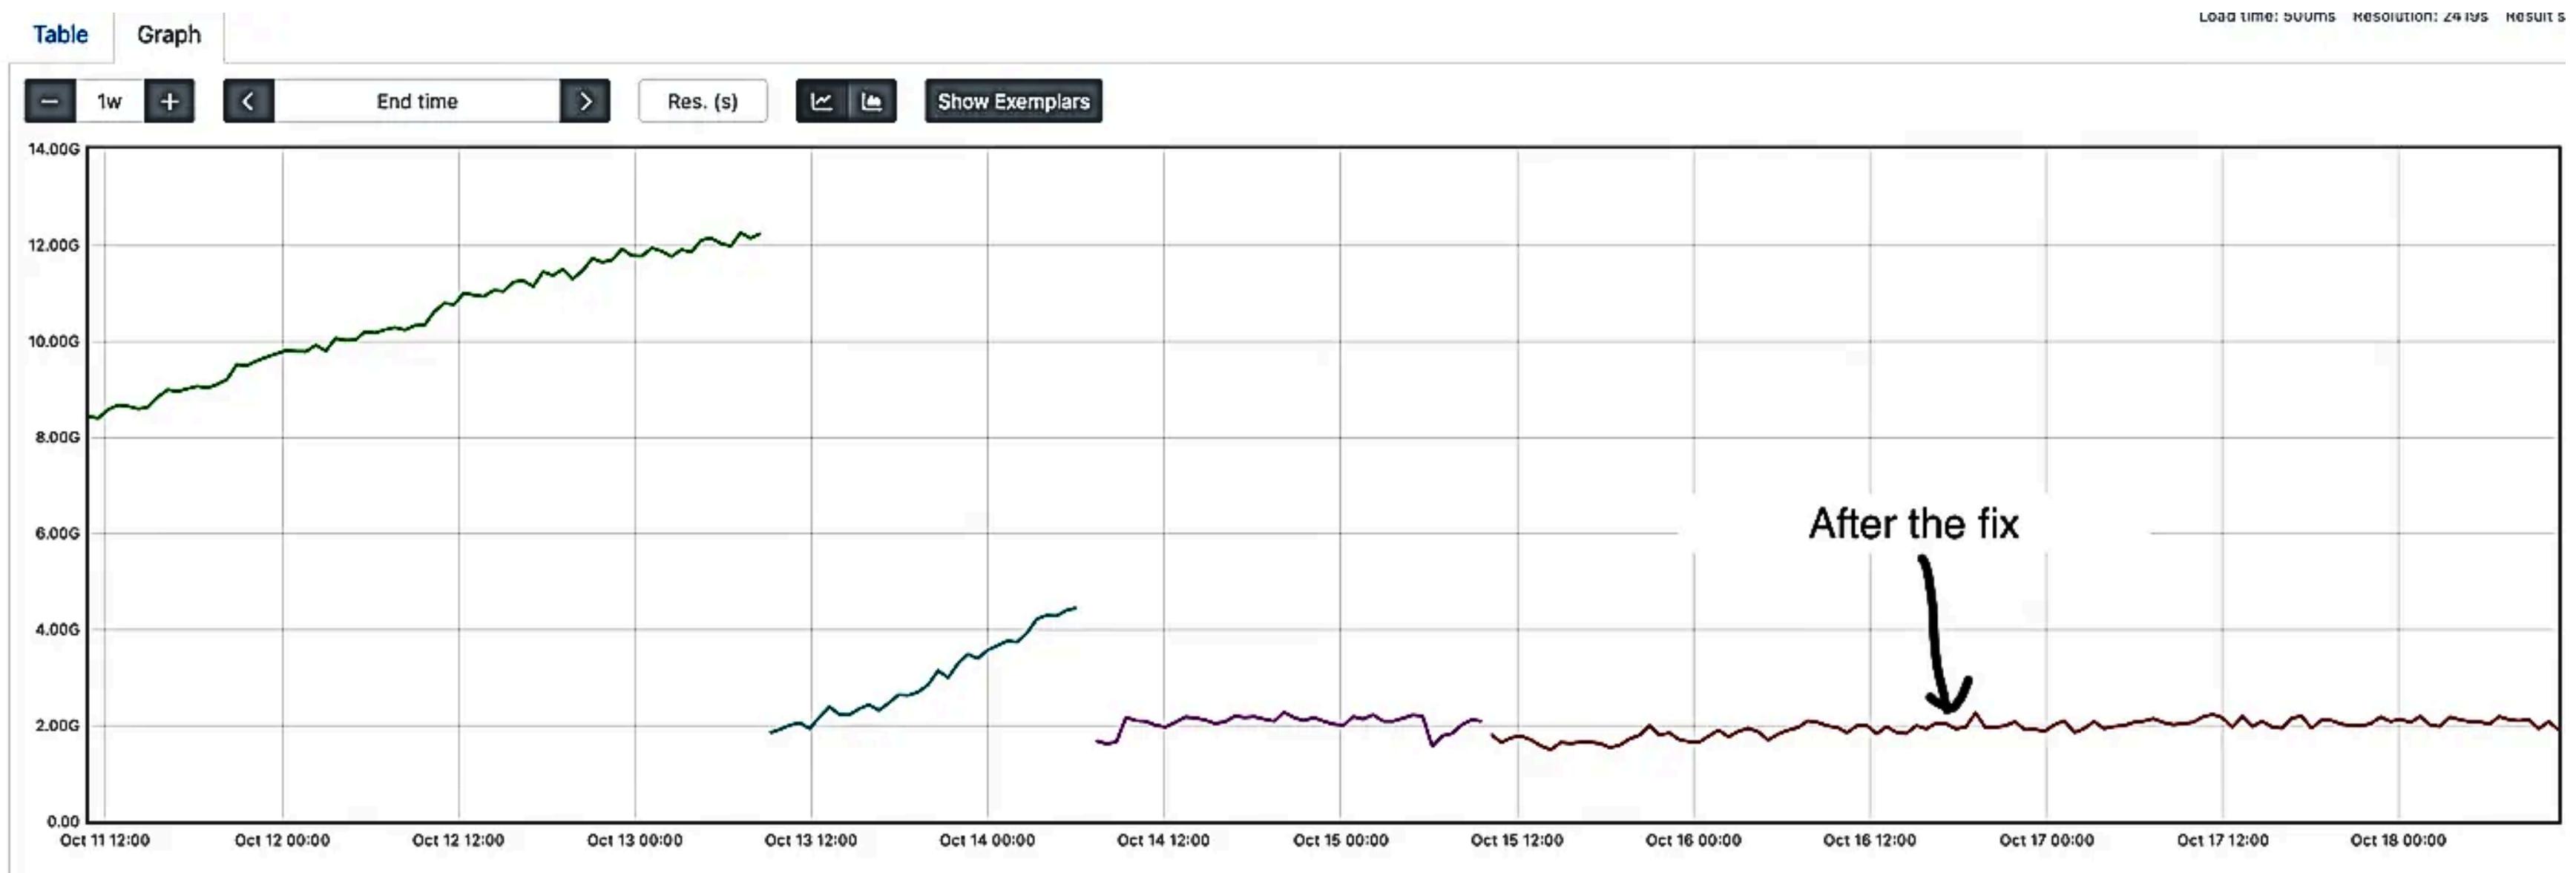
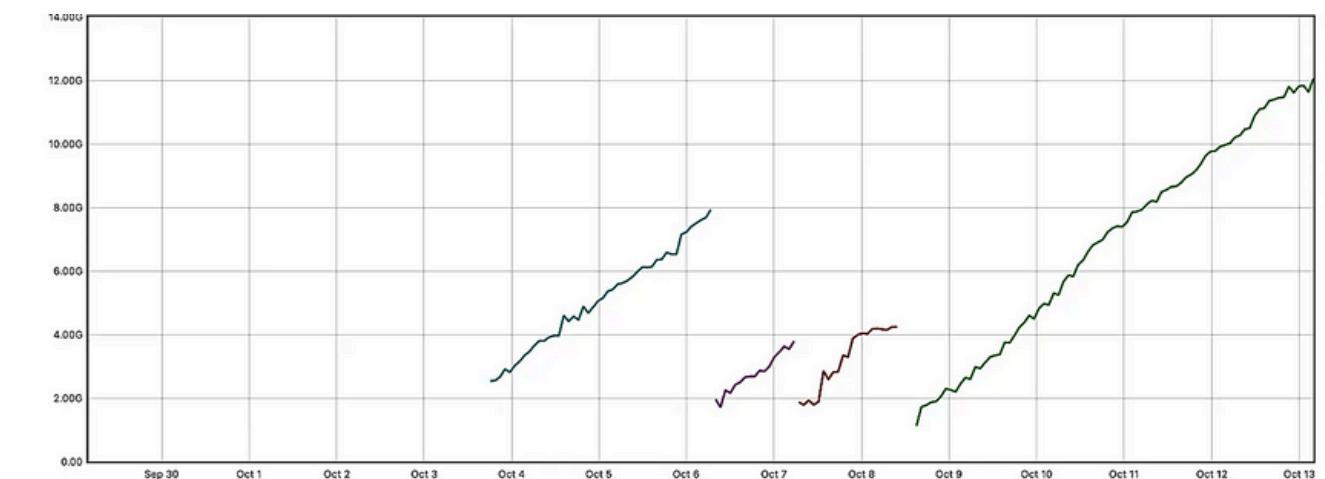
Assignees  
nickitat

Labels  
can be tested pr-



<https://github.com/ClickHouse/ClickHouse/pull/88617>

Memory leak =>



# Network errors in single pod Clickhouse logs

## Error type 1:

Response: {"message": "Invalid Argument: 6  
:Internal Error: 5  
:clickhouse: Code: 279. DB::NetException: **All connection tries failed.** Log:  
Code: 210. DB::NetException: I/O error: **Broken pipe**, while writing to socket (127.0.0.1:40882 -> 127.0.0.1:9000).  
**(NETWORK\_ERROR)** (version 24.3.5.46 (official build))  
Code: 210. DB::NetException: I/O error: Broken pipe, while writing to socket (**127.0.0.1:40896 -> 127.0.0.1:9000**).  
**(NETWORK\_ERROR)** (version 24.3.5.46 (official build))



github.com/ClickHouse/ClickHouse/issues/64954#issuecomment-2365285608

Closed

The NETWORK\_ERROR occurs frequently #64954



alexey-milovidov on Sep 22, 2024

The diagnostic in the logs is correct. The error messages are correct.



4



alexey-milovidov closed this as completed on Sep 22, 2024

You have a flaky network,  
go run clickhouse in a stable setup please!!



# Network errors in single pod Clickhouse logs

## Error type 2:

2025.03.05 04:53:13.804881 [ 1153 ] {c349a22c-7450-4f49-85d7-17e430386702} <Information>

**HedgedConnectionsFactory**: Connection failed at try №1, reason: Code: 210. DB::NetException: I/O error:  
Broken pipe, while writing to socket (**127.0.0.1:47536 -> 127.0.0.1:9000**). (**NETWORK\_ERROR**) (version  
24.8.8.17 (official build))

## How Hedged Requests Work:

1. ClickHouse starts by sending a query to one replica.
2. If the response is delayed beyond a certain threshold, it sends the same query to another replica.
3. The first response to return is used, and the other request(s) are discarded.

# Network errors in single pod Clickhouse logs

## Error Type 3:

stack\_trace: 0. DB::Exception::Exception(DB::Exception::MessageMasked&&, int, bool) @ 0x00000000d17749b  
1. DB::NetException::NetException<String&>(int, FormatStringHelperImpl<std::type\_identity<String&>::type>, String&) @ 0x0000000011c35703  
2. **PoolWithFailoverBase<DB::IConnectionPool>**::getMany(unsigned long, unsigned long, unsigned long, unsigned long, bool, bool,

Connection Pools related to Distributed query processing exhausted

# Found connection configs, reviewed and updated

Clickhouse config	Default	Updated values	Description
<b>tcp_keep_alive_timeout</b>	290s	1s	The time in seconds the connection needs to remain idle before TCP starts sending keepalive probes
<b>use_hedged_requests</b>	1	0	Use hedged requests for distributed queries
<b>connections_with_failover_max_tries</b>	3	10	The maximum number of connection attempts with each replica for the Distributed table engine.
<b>distributed_connections_pool_size</b>	1024	4096	The maximum number of connection attempts with each replica for the Distributed table engine.
<b>max_distributed_connections</b>	1024	64	The maximum number of connections for distributed processing of one query (should be greater than max_threads).

Caution: Values depend on the nature of **your** environment

# Questions?



## Thank you!