



KubeCon



CloudNativeCon

India 2025

Bridging Data & ML ecosystems: A cloud Native Approach using Kubeflow

Johnu George

Shivji Kumar Jha

About the Speakers



Johnu George
Technical Director
Nutanix AI

- Kubeflow Steering Committee Member
- Kubeflow Training & AutoML WG chair
- MLCommons Storage Chair



Shivji Kumar Jha
Staff Engineer
Data Platforms at Nutanix

- Software Engineer: Distributed systems / databases
- Contributed code to MySQL, Pulsar, Clickhouse
- Excited about Open-Source Software & Communities
- Past Talks: <https://github.com/shiv4289/shiv-tech-talks/>

Table of Contents



Kubeflow introduction



ML training architecture



Data challenges



Solutions

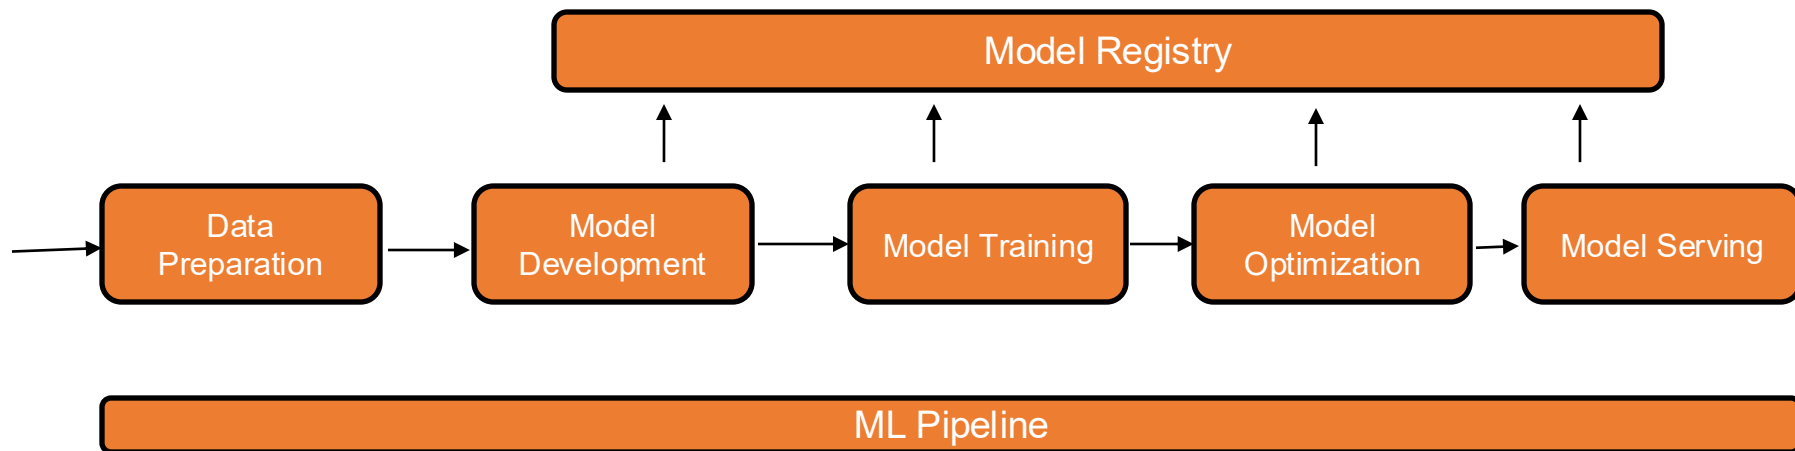


Putting it all together

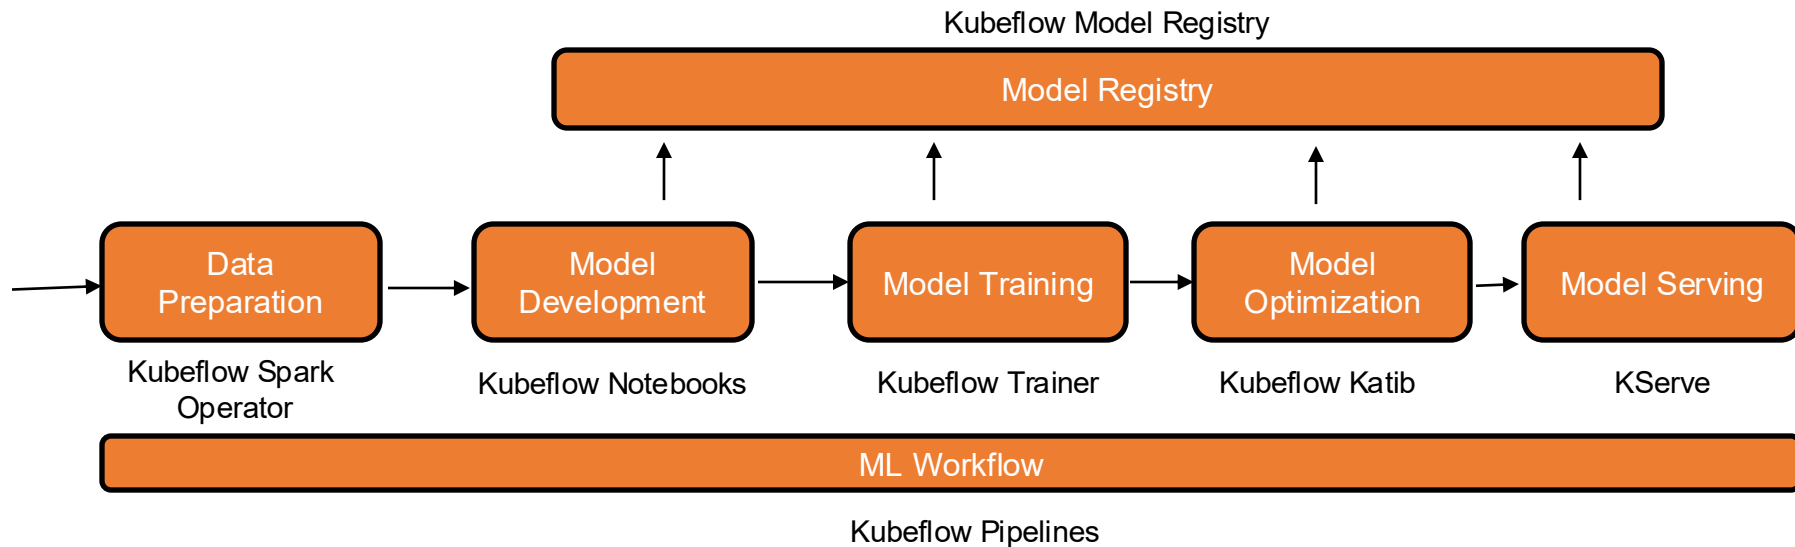


Demo code!

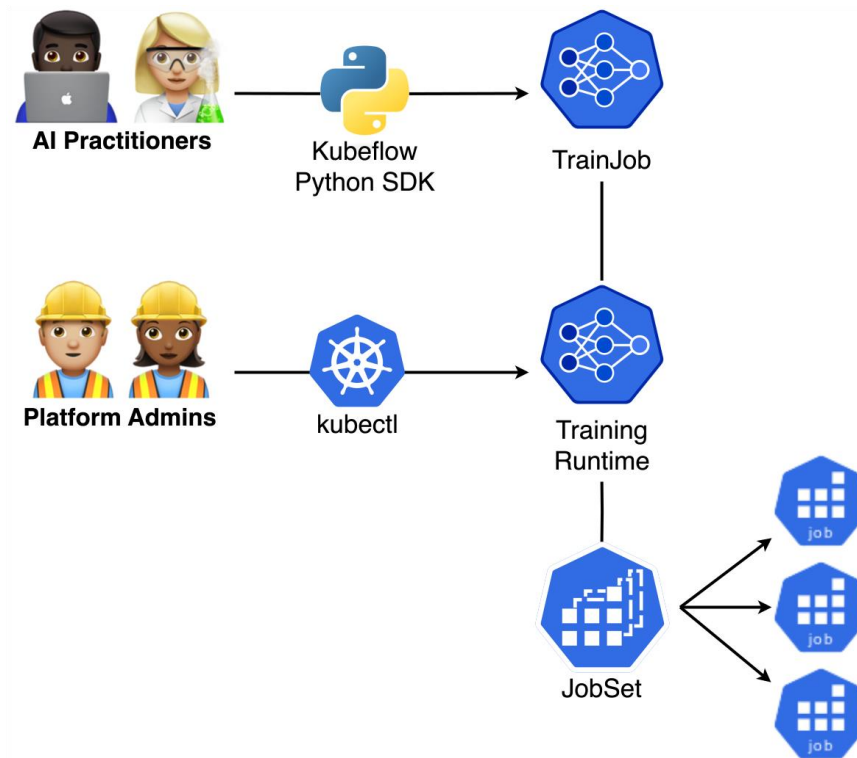
MLOps Pipeline



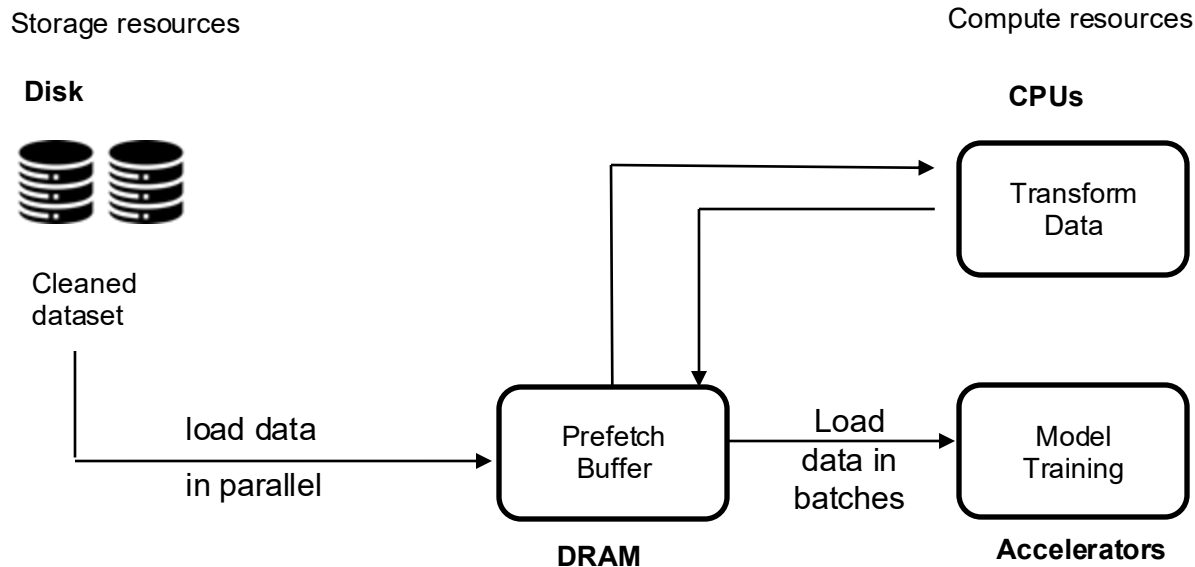
Kubeflow Platform



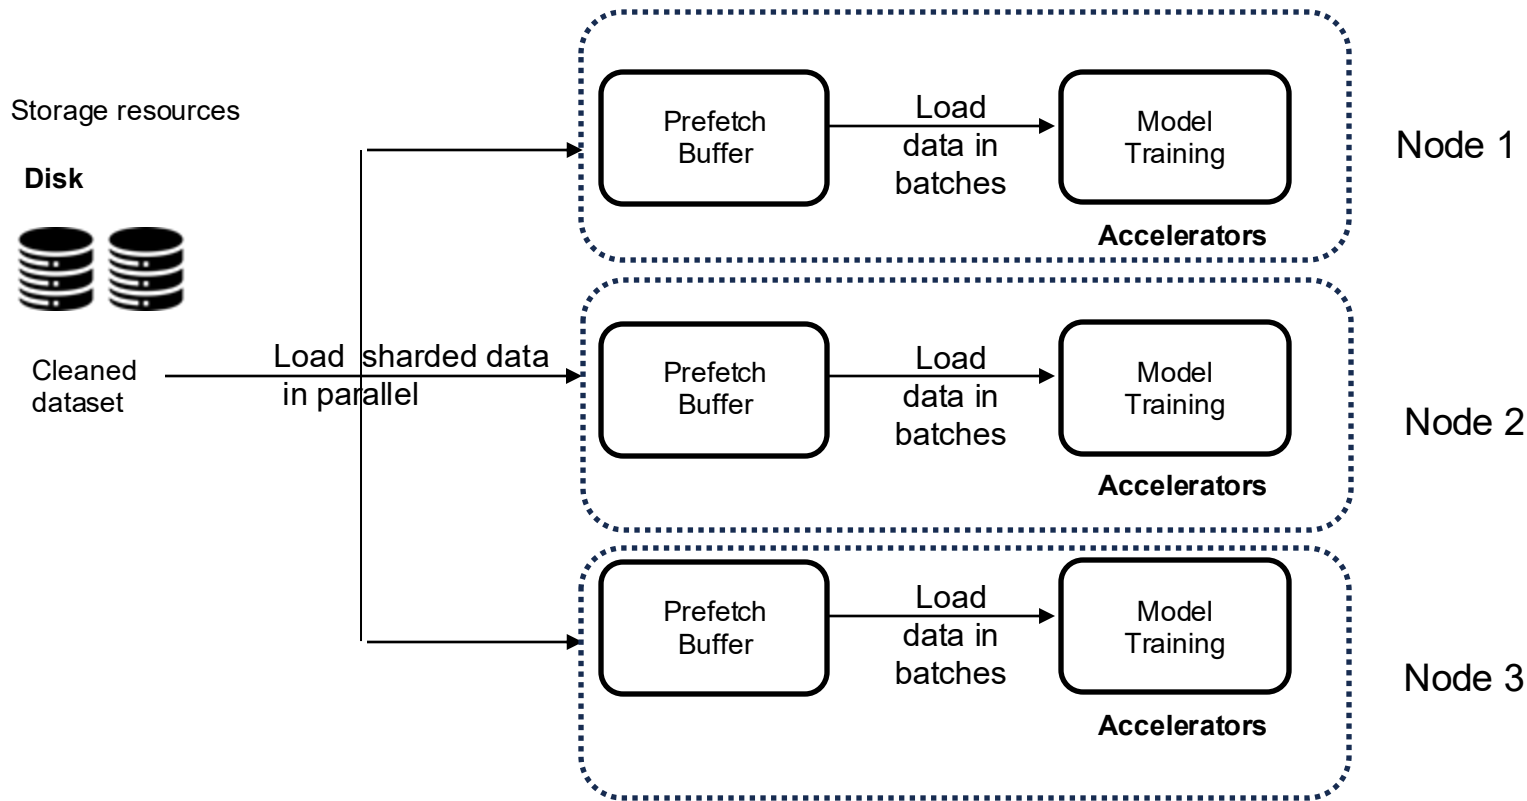
Kubeflow Trainer



Single Node ML training



Multi-Node ML training (DDP)



Challenges in ML stack

Async storage
access

- Efficient storage for large datasets
 - Larger dataset, better the model
- GPU under utilization
 - High data throughput requirement to keep GPU busy
- Sharding data across workers
 - Unique subset of data for every worker
- Consistent shuffling of data across epochs
 - Restream data in every epoch
 - Network bandwidth wastage
- Supporting heterogeneous training jobs on same dataset
 - I/O wastage due to static data sharding pattern

Repeat for n epochs

for epoch in range(n_epochs):

for i, data in enumerate(training_loader):

Every data instance is an input + label pair
inputs, labels = data

Zero your gradients for every batch!
optimizer.zero_grad()

Make predictions for this batch
outputs = model(inputs)

Compute the loss and its gradients
loss = loss_fn(outputs, labels)
loss.backward()

Adjust learning weights
optimizer.step()

Next-gen data stack



Apache Parquet



Large data set needs efficient storage



Parquet is a **compressed, columnar storage format**



Lower storage & IO cost.

Easy Access to Parquet Files

Apache Iceberg



S3 means - Infinite Storage & good HA



Open Table format (SQL & ACID)




Lots of integrations -> access patterns



But... slow access



And high transfer costs



SUBWAY

analogy

Grab prepped ingredients & cook fast

WAREHOUSE



Bulk Storage
(Frozen Veggies, meats,
sausage)

WAREHOUSE

CACHE


TRAINERS

Pull needed, warm,
clean, sort; ready to pick




analogy

WAREHOUSE



ICEBERG



CACHE



PYTORCH

TRAINERS

Introducing a Cache for trainers



ICEBERG

Store huge datasets
in a cost-effective &
open format



CACHE

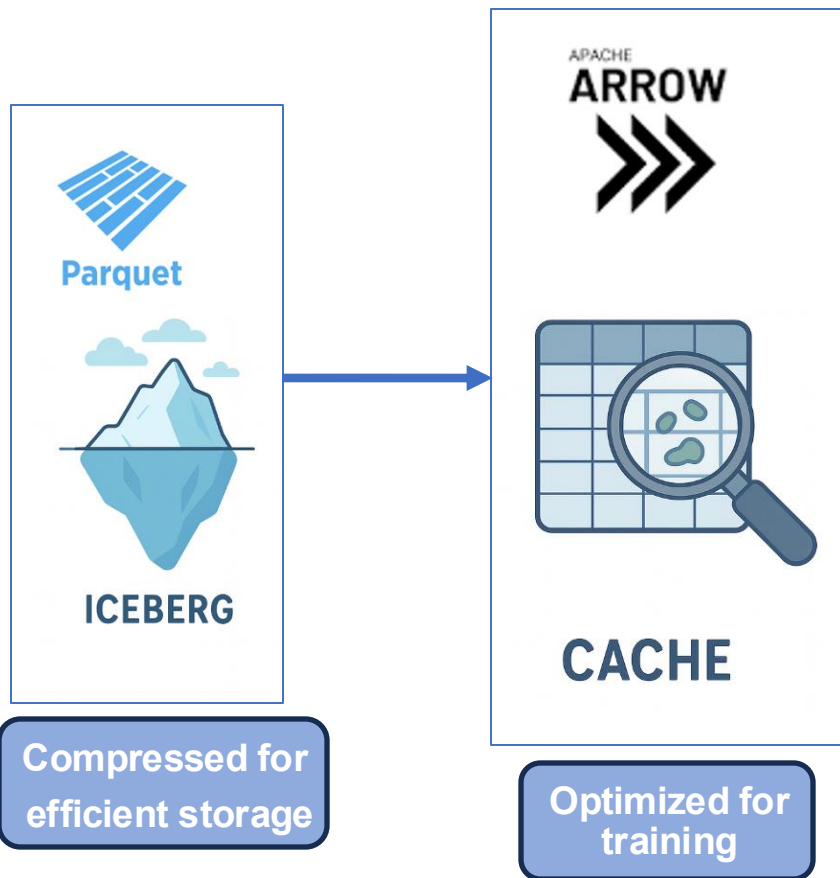
For easy access of
prepped data to
training workers



PYTORCH

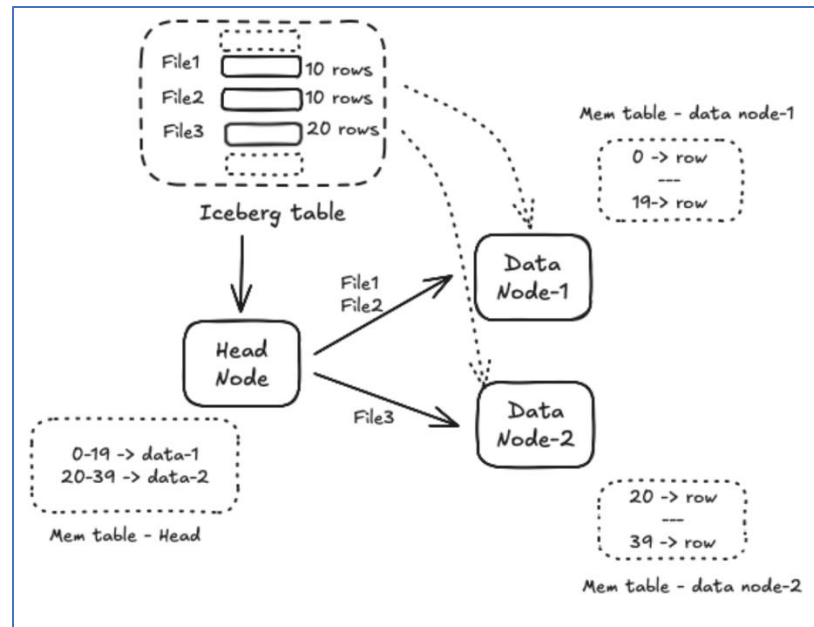

Distributed
GPU-enabled
ML training workers

Cache Implementation

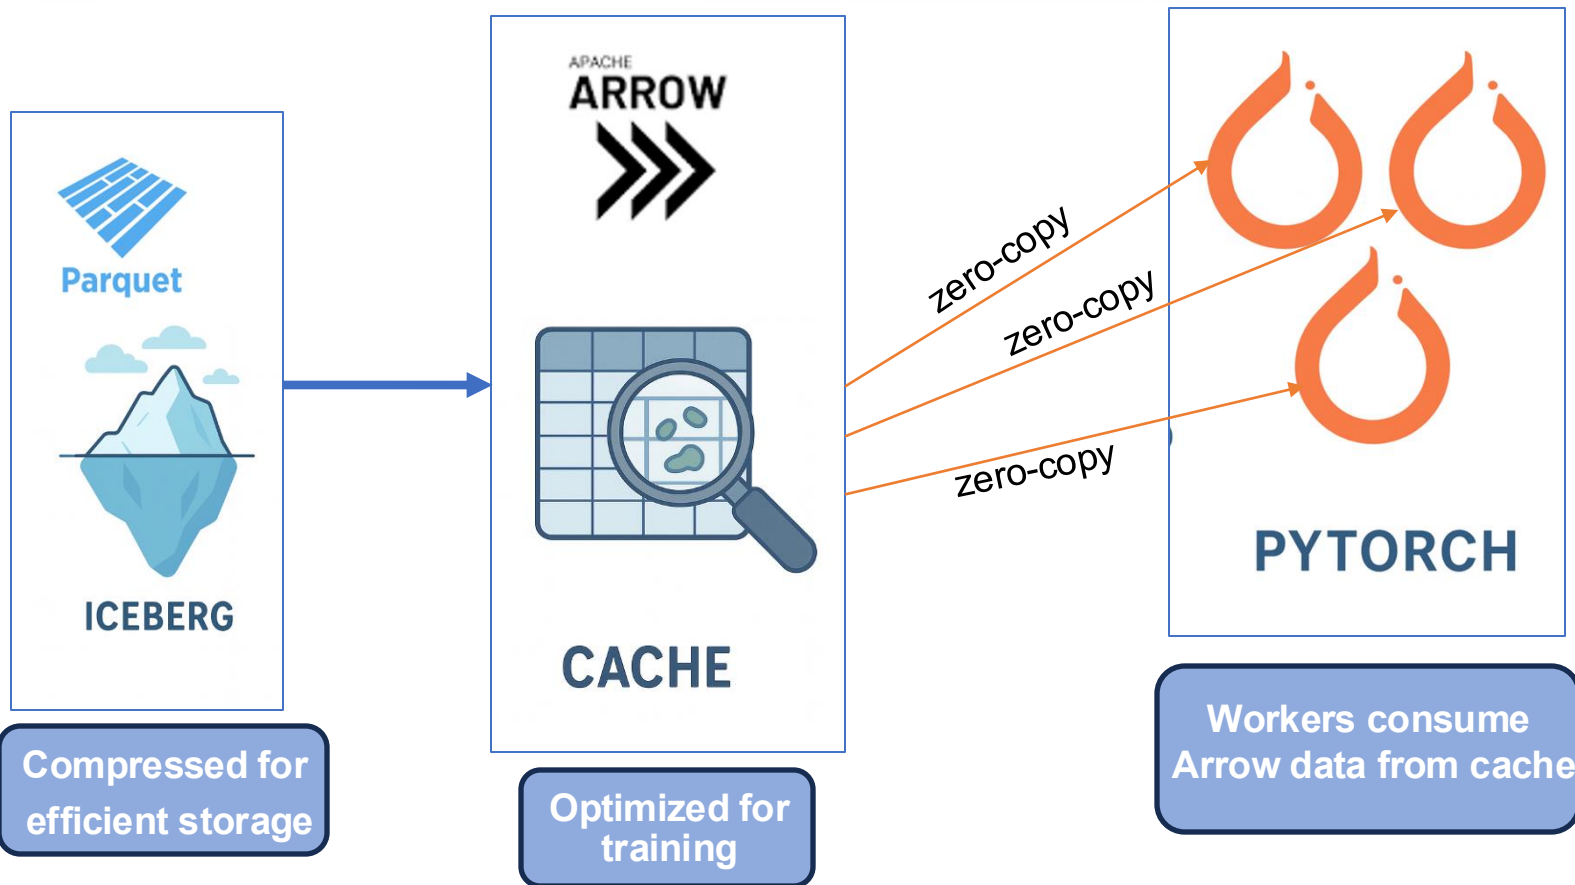


- Cache Data is kept in Arrow format
- Efficient in-memory storage & transport
- Cache nodes convert Parquet from (Iceberg → Arrow RecordBatches)

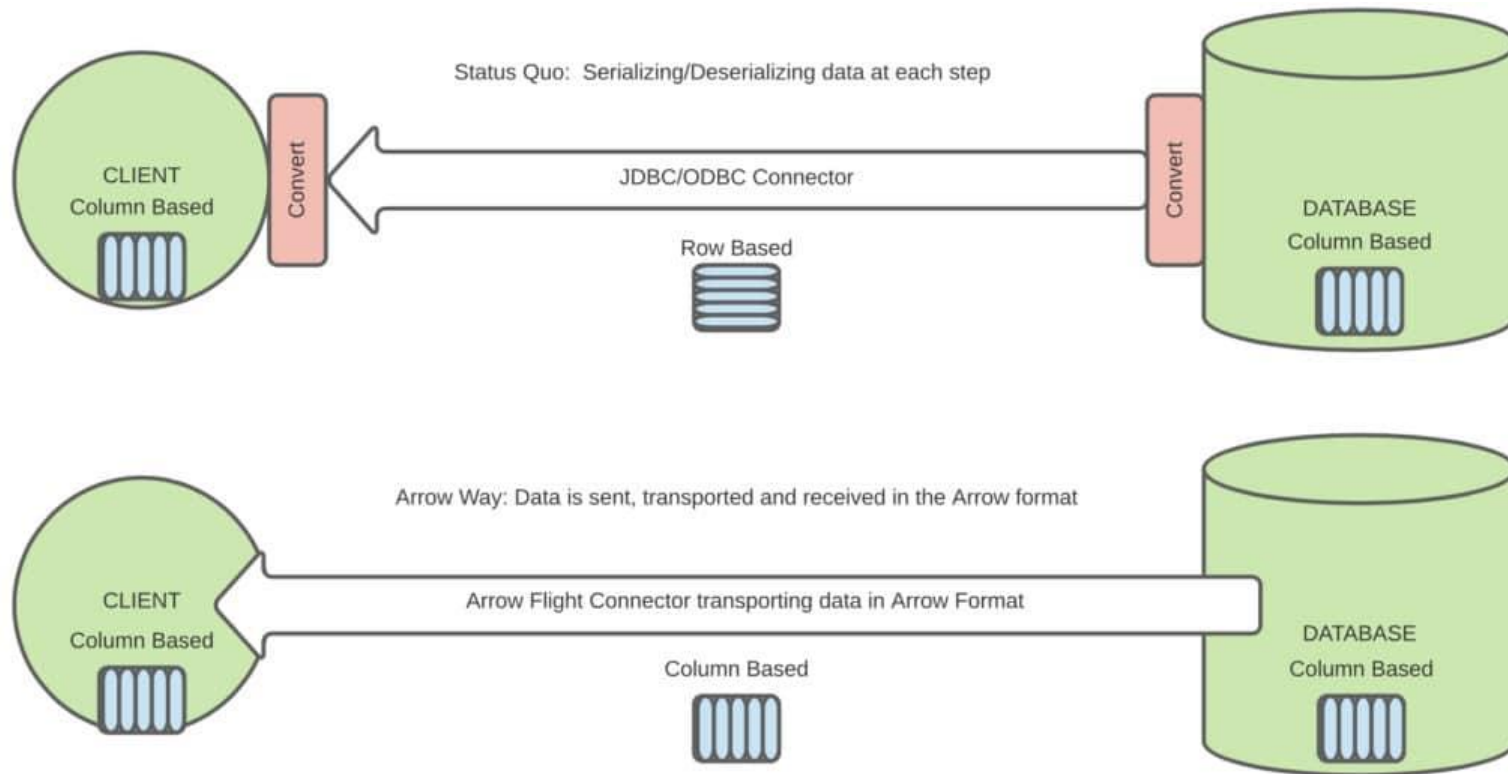
Moving Data from Iceberg to Cache



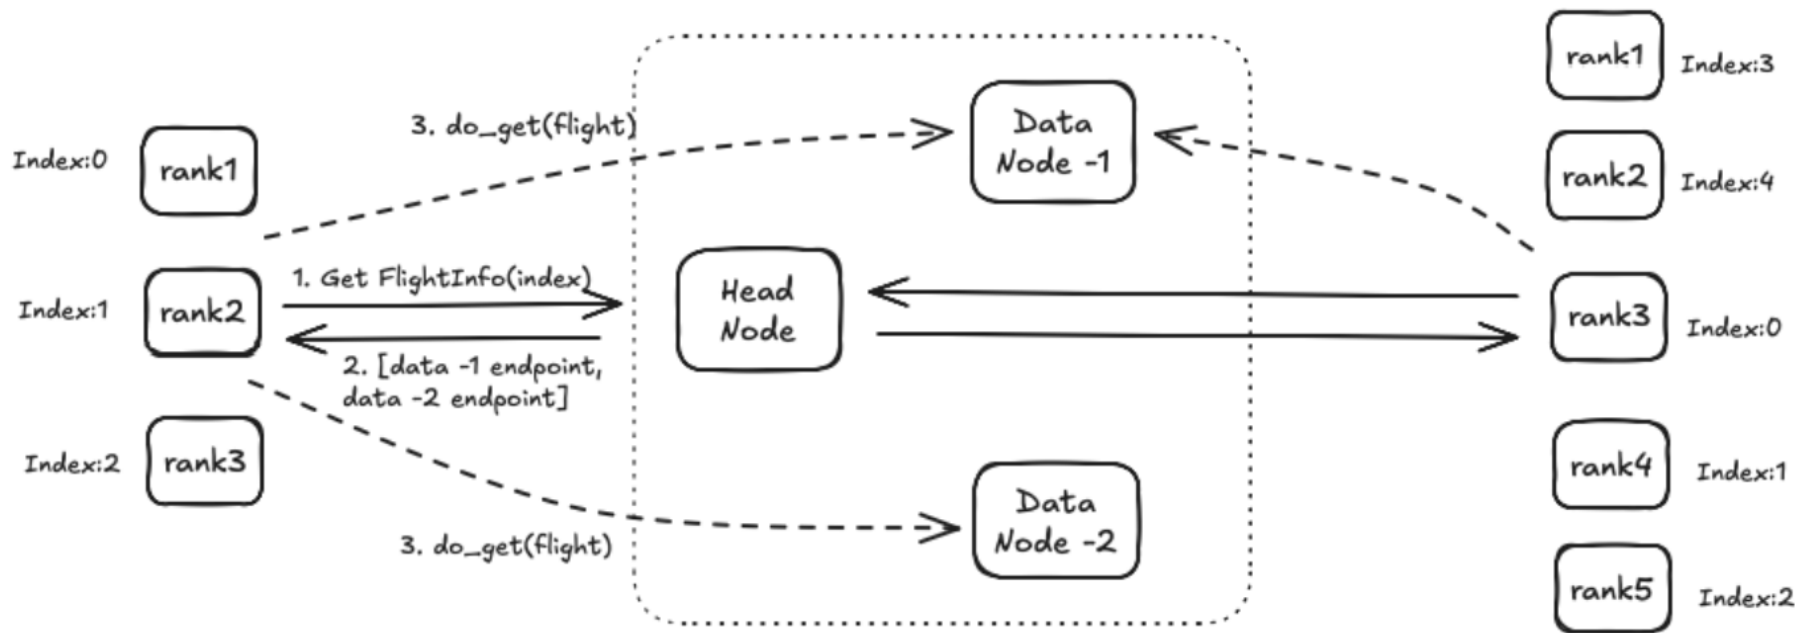
Fetching data from Cache



Arrow Flight: zero-copy transport



A Cache built for training workers



 TrainJob-1



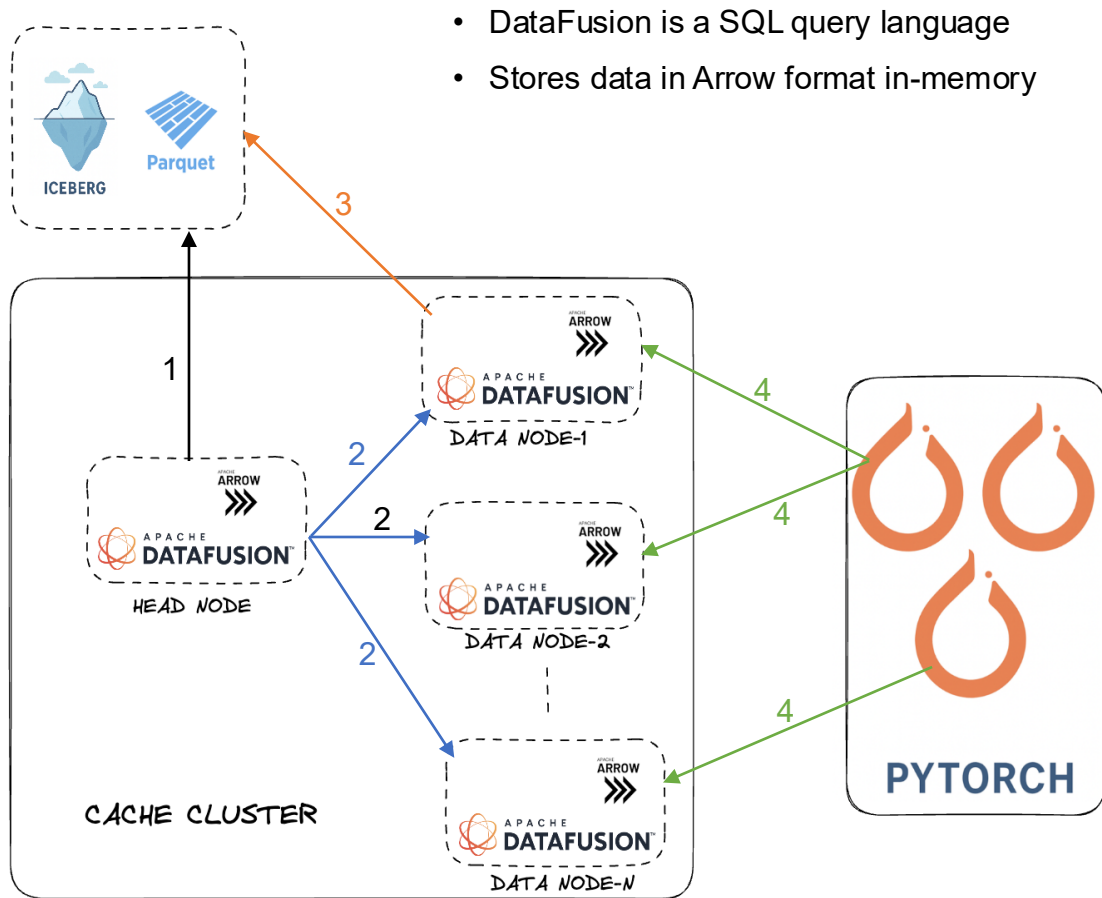
CACHE

TrainJob-2 

Putting it all together

DataFusion

1. Queries Metadata from Iceberg
2. Head node assigns slices to data nodes.
3. Data nodes fetch, store data (arrow)
4. PyTorch workers fetch required data



Revisiting the challenges

1. Efficient storage for large datasets

- Larger dataset, better the model

Apache Parquet

2. GPU under utilization

- High data throughput requirement to keep GPU busy

3. Consistent shuffling of data across epochs

- Restream data in every epoch
- Network bandwidth wastage

Arrow Flight

Custom Infra for
ML training

4. Sharding data across workers

- Unique subset of data for every worker

indexed data storage

5. Supporting heterogeneous training jobs on same dataset

- I/O wastage due to static data sharding pattern

Dynamic assignment
based on #workers

Let's write some code with Kubeflow SDK!

“Hello cache!”: with new data stack



KubeCon



CloudNativeCon

India 2025

```
client.train(
runtime_ref="torch-distributed-with-cache",
trainer=CustomTrainer(
    func=train_func,
    num_nodes=3,
    resources_per_node={
        "cpu": 3,
        "memory": "6Gi",
        "nvidia.com/gpu": 1,
    },
),
initializer =
    Initializer(dataset=ArrowCacheDatasetInitializer(
        cluster_size="3",
        metadata_loc="s3a://testbucket/feature_demo/data/yelp_review_v2/metadata/1234.metadata.json",
        table_name="yelp_data_v2",
        schema_name="feature_demo",
        features:[],
        filter:"",
    )),
Runtime=torch.runtime
)
```


Join us!



Huge shoutout to Akshay Chitneni, Andrey Velichkevich, Rasik Pandey

Design: [Google Doc](#)

Issue Tracker: <https://github.com/kubeflow/trainer/issues/2655>

Proposal: <https://github.com/kubeflow/community/pull/864>

Implementation : <https://github.com/kubeflow/trainer/pull/2755>

Call for contributors



KubeCon



CloudNativeCon

India 2025

Questions?

Get in touch:

Johnu George - <https://www.linkedin.com/in/johnu-george-83036610/>

Shivji Kumar Jha - <https://www.linkedin.com/in/shivjiijha/>